

1 Bridging the Gap: Precise Static Analysis of WebAssembly in 2 a JavaScript World 3

4 ANONYMOUS AUTHOR(S)
5

6 WebAssembly is a low-level bytecode that runs alongside JavaScript to enable near-native performance on the
7 web. Static analysis of WebAssembly is challenging, particularly in the case of call graph analysis because the
8 dispatch of function calls is done using non-constant values with complex data-flow. Existing state-of-the-art
9 analysis tools do not perform a pointer and value analysis which is essential for precise call graph analysis.
10 Additionally, they analyze the WebAssembly binary in isolation and make worst-case assumptions about a
11 binary's JavaScript client, which leads to loss of precision in a call graph analysis.

12 We address this gap with WASMBRIDGE, the first multi-language static analysis tool that integrates in-
13 teroperation between WebAssembly and JavaScript. WASMBRIDGE introduces a refinement type system for
14 Wasm that distinguishes numeric values from pointers and types WebAssembly memory to recover types for
15 marshalled JavaScript data. We prove the type system sound and show that incorporating pointer information
16 substantially improves the precision of call graph construction. Across a corpus of real-world Wasm applica-
17 tions, WASMBRIDGE reduces the number of estimated call targets at 35% of indirect call sites—by up to 17% per
18 site—yielding an overall call-graph edge reduction of 5% compared to state-of-the-art Wasm analyses. We
19 further demonstrate the utility of our approach through dead-code elimination, identifying up to 20% more
20 dead functions than an industry tool. WASMBRIDGE provides an extensible foundation for whole-program,
21 multi-language analysis of JavaScript-WebAssembly applications, enabling future integration with advanced
JavaScript value and call graph analyses.

22 CCS Concepts: • Software and its engineering → Automated static analysis.

23 Additional Key Words and Phrases: WebAssembly, JavaScript, Interoperation, Multi-Language Static Analysis,
24 Static Analysis, Constraints, Debloating, Dead Code Elimination

25 1 INTRODUCTION

26 Program analysis is essential to understanding the functionality and behavior of a binary program
27 when its source code is unavailable. It has been used extensively to detect malware and vulnerabili-
28 ties in binary programs [23, 35], to reverse engineer binaries [24] and to optimize binaries outside the
29 compiler toolchain [2]. However, program analysis of binaries is also very challenging since a sub-
30 stantial amount of information is lost during compilation. Compilers erase types, variable and func-
31 tion names, and all high-level structure present in the source to create the most performant binary
32 possible. This loss of information, coupled with the under-specified semantics of low-level languages
33 and large instruction sets, makes writing sound and precise analyses for binaries difficult [3].

34 The problem of binary analysis is now also on the Web! With the introduction of WebAssembly
35 (Wasm) in 2017 [20], compilers have been able to compile languages like C, C++, Rust, etc. to
36 WebAssembly, a portable, low-level bytecode format that has been designed for computationally
37 intensive tasks in the browser. These WebAssembly binaries interoperate with a JavaScript client
38 on browsers and in Node.js applications. These applications don't include the source code of the
39 binary and so, analysis writers for these applications must reckon with the hardships of binary
40 analysis. Some of these hardships are mitigated by WebAssembly's formally specified semantics [40],
41 structured control flow and low-level value types. Prior work [28] succinctly captures the specific
42 challenges of WebAssembly static analysis with the example of call graph analysis. In WebAssembly,
43 indirect calls are mediated through a table containing a list of functions, where an integer index is
44 used to call a specific function from the table. Standard static analysis techniques dictate that a
45 value analysis be used to determine the index used in indirect calls. However, Lehmann et al. [28]
46 identify several unique challenges to WebAssembly call graph analysis, which they empirically
47 validate over a dataset of 8,461 WebAssembly binaries [21]. They find that the dispatch of function
48

calls is done using non-constant values with complex data-flow in nearly all cases, complicating the task at hand. Furthermore, since WebAssembly commonly interoperates with JavaScript, aspects of this interoperation pose further challenge to call graph analysis. For example, JavaScript can mutate exported WebAssembly function tables. Thus, to precisely determine the value of an index used for indirect calls, a WebAssembly static analysis would have to have both a sophisticated value and pointer analysis, and an analysis of the JavaScript client.

However, few WebAssembly static analysis tools have risen to the task. Tools like WASSAIL[36] either limit themselves to a constant index analysis, which is not found to occur in practice [28], while tools like STURDY[22] use abstract interpretation to track the data flow of abstract integers inter-procedurally, but do not perform pointer analysis beyond that of constant memory addresses. Additionally, to the best of our knowledge, no Wasm static analysis tool supports passing in values from JavaScript. Current state-of-the-art tools either analyze WebAssembly in isolation or make worst-case assumptions about JavaScript client behavior, such as that a client always mutates the WebAssembly function table. Adding support for interoperation is non-trivial. While JavaScript Numbers and BigInts are passed to WebAssembly functions directly, other datatypes (strings, arrays, objects) are marshalled into WebAssembly memory and passed to Wasm functions as pointers to memory. No static analysis for WebAssembly performs pointer analysis at all.

In this paper, we address the central problem: the lack of support for interoperation with JavaScript in the WebAssembly static analysis literature. We perform a multi-language analysis to support passing in JavaScript values into a WebAssembly static analysis. Current value analyses cannot perform a multi-language analysis because of a lack of representation of pointers in the WebAssembly type system. WebAssembly has a low-level type system of 32- and 64-bit integer types (i32, i64) and 32- and 64-bit floating point types (f32, f64) and WebAssembly memory is untyped. We design a refinement type system that refines i32 types to either a singleton pointer type $\text{ptr}(l, n)$ or a singleton number $\text{num}(n)$ — where $\text{ptr}(l, n)$ refers to a pointer to location l at offset n . We also type the WebAssembly memory, which lets us recover JavaScript values marshalled into the memory. We prove this refinement type system to be type-safe and show that the addition of a pointer analysis improves the precision of call graph analyses over standalone WebAssembly binaries. We improve precision at an average of 35% of indirect call sites, reduce the number of edges in the call graph by up to 5% compared to other Wasm static analysis tools, and discover a small number of call sites to be monomorphic, which enables additional compiler optimizations.

Along with our refinement type system, we also implement an analysis of the JavaScript code responsible for the marshalling and unmarshalling of JavaScript values to WebAssembly. While this does not improve precision of the WebAssembly analysis, it makes our tool extensible to other JavaScript client analyses. Specifically, to do a precise *open-world analysis* of a JavaScript-WebAssembly application, a tool that performs a value and call graph analysis of the JavaScript client can supply WASMBRIDGE with JavaScript values passed to the relevant subset of WebAssembly-exported functions. We demonstrate the precision of such a multi-language analysis by performing dead-code elimination to specialize a WebAssembly binary to a JavaScript client. Here we are able to find up to 20% more functions to be dead than an industry tool, WASM-METADCE[41] does. To the best of our knowledge, there is no off-the-shelf sophisticated value and call graph analysis for entire JavaScript applications. Thalakottur et al. [38] find that WebAssembly is often depended upon in JavaScript applications through complex indirect dependencies, which further complicates the analysis. Therefore, we leave such analyses to future work.

We summarize the contributions of this work as follows:

- (1) We present the first multi-language static analysis of WebAssembly and JavaScript, WASMBRIDGE, that analyzes JavaScript wrappers and performs an intra-procedural pointer and value analysis over WebAssembly binaries.

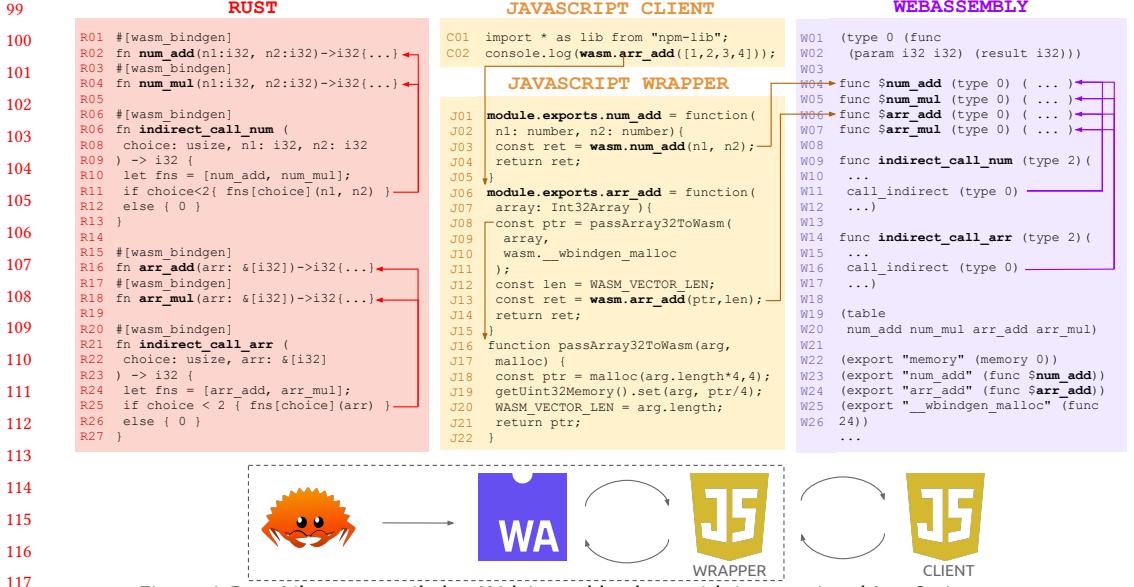


Fig. 1. A Rust Library compiled to WebAssembly along with its associated JavaScript wrapper.

- (2) We develop a refinement type system for WebAssembly that refines the WebAssembly value type i32 to distinguish between numbers and pointers to WebAssembly memory.
- (3) We show that our approach, WASMBRIDGE, leads to a reduction in the number of estimated targets at an average of 35.63% of indirect call sites, with an average reduction of upto 17% at those call sites.
- (4) We evaluate WASMBRIDGE against current state-of-the-art static analyses for WebAssembly and show an edge reduction of 5% when compared against those tools.
- (5) We estimate the impact of the increased precision of our analysis on dead-code elimination and find up to 20% more dead functions than a state-of-the-art industry tool, WASM-METADCE.

The remainder of this paper is organized as follows. Section 2 provides background on the interoperation between WebAssembly and JavaScript and introduces the intuition behind our approach. Section 3, 4 and 5 delve into the specifics of our approach. We describe our evaluation against real-world subjects in Section 6. Section 7 and 8 discuss related and future work respectively. WASMBRIDGE is anonymously available at <https://anonymous.4open.science/r/wasmbridge-eval-2D7B/>.

2 BACKGROUND AND MAIN IDEAS

2.1 Interoperation Through a Wrapper

We explain the interoperation between WebAssembly and JavaScript using the example shown in Figure 1, where we show a Rust library being compiled to WebAssembly and its associated compiler-generated JavaScript wrapper file. The library functions in the wrapper file are then called from a JavaScript client. We use different colors for each language for clarity.

2.1.1 The Library Developer Perspective. Let us consider the Rust library shown in Figure 1. The library consists of operations over numbers and arrays. `num_add` adds two numbers, `num_mul` multiplies two numbers, `arr_add` returns the sum of the elements of an array and `arr_mul` returns the result of multiplying all the elements of an array. The library also has two functions, `indirect_call_num` and `indirect_call_arr` that are used to indirectly call one of the number or array operations based on some user-specified choice. A Rust library developer who wants to compile

148 this library to WebAssembly annotates the functions in the library that they want to expose to a
 149 JavaScript client with a special tag, `#[wasm_bindgen]`¹. We refer to functions that are exposed to
 150 the JavaScript client as *exported functions*. The reader can see in Figure 1 that all the library functions
 151 are exported. The JavaScript client is meant to call these exported functions with JavaScript objects.
 152 The translation from **JavaScript** values to **Rust** values is handled by `wasm_bindgen` and is not
 153 something that the Rust library developer has to worry about. For exported Rust functions, the
 154 compiler produces a WebAssembly binary along with a JavaScript wrapper file. For libraries being
 155 used by Node.js applications, these files (without the Rust source) are bundled in an NPM² package.
 156 A JavaScript client imports this NPM package to use the library from JavaScript. In Figure 1, the
 157 Rust library is bundled in the NPM package "npm-lib".

158 2.1.2 *The JavaScript Client Perspective.* A JavaScript client imports the library NPM package as
 159 an object `lib` (line C01). It then calls an exported function with JavaScript objects. On line C02,
 160 we see a call to `arr_add` with an array. The client does not translate this **JavaScript** object into a
 161 **WebAssembly** value. In fact, they do not interact with WebAssembly directly at all!

162 2.1.3 *The JavaScript Wrapper.* In real-world applications, JavaScript and WebAssembly interoperate
 163 through a wrapper: a JavaScript file generated by the compiler alongside the binary. This wrapper
 164 is used to *lower* a JavaScript object to WebAssembly before a call to a WebAssembly function and
 165 to *lift* a WebAssembly value to a JavaScript object after a call to a WebAssembly function. The
 166 code that does the lifting and lowering is often referred to as glue code. A WebAssembly module
 167 exposes certain functions to JavaScript through exports (lines W22-W26). The wrapper file contains
 168 JavaScript functions that have glue code around calls to exported WebAssembly functions.

169 To understand the glue code that mediates the interoperation between WebAssembly and
 170 JavaScript, let us go back to the example in Figure 1. In function `num_add` (line J01), we see
 171 no glue code before the call to `num_add`. This is because the JavaScript engine internally translates
 172 a JavaScript `number` to a WebAssembly `i32` value. However, we see that in function `arr_add` (line
 173 J06), the `Int32Array` object passed into the function is passed to function `passArray32ToWasm` along
 174 with another exported WebAssembly function, `__wbindgen_malloc`. In the body of the function,
 175 we first see a call to `__wbindgen_malloc` with the size of the array in bytes. `__wbindgen_malloc`
 176 allocates a block of WebAssembly memory of the required size and returns a pointer to this block.
 177 Then a call to the function `getUint32Memory` accesses the WebAssembly memory³ and updates
 178 the block at the pointer with the contents of the array. The pointer and the array length, both We-
 179 bAssembly `i32` values, are then passed in as arguments to `arr_add` (line J13). Since the WebAssembly
 180 type system is fairly low-level, this is an expected pattern. JavaScript objects that cannot be passed
 181 to Wasm directly are stored in memory and the pointer to memory is passed to Wasm instead.

183 2.2 Call-Graph Analysis Over the WebAssembly Binary

184 A call-graph analysis over the source code of the Rust library will resolve the targets for the indirect
 185 calls at R11 and R25 as `num_add`, `num_mul` and `arr_add`, `arr_mul` respectively. Unfortunately, a
 186 call-graph analysis over the corresponding WebAssembly binary produced by the compiler is not
 187 as precise. Indirect calls in WebAssembly occur using a `call_indirect` instruction and are mediated
 188 through a WebAssembly table that contains a list of functions (line W19). WebAssembly is a
 189 stack machine and instructions push and pop values from the stack. At runtime, the `call_indirect`
 190 instruction pops an `i32` value from the stack which it uses to index into the table, thus determining

191 ¹wasm-bindgen is a popular tool that enables Rust-JavaScript interoperation by generating bindings and glue code for
 192 JavaScript clients looking to use a Rust library through WebAssembly. We discuss this glue code in detail shortly.

193 ²NPM is the canonical package manager for Node.js applications.

194 ³Note that the WebAssembly memory is also exported to JavaScript in W22.

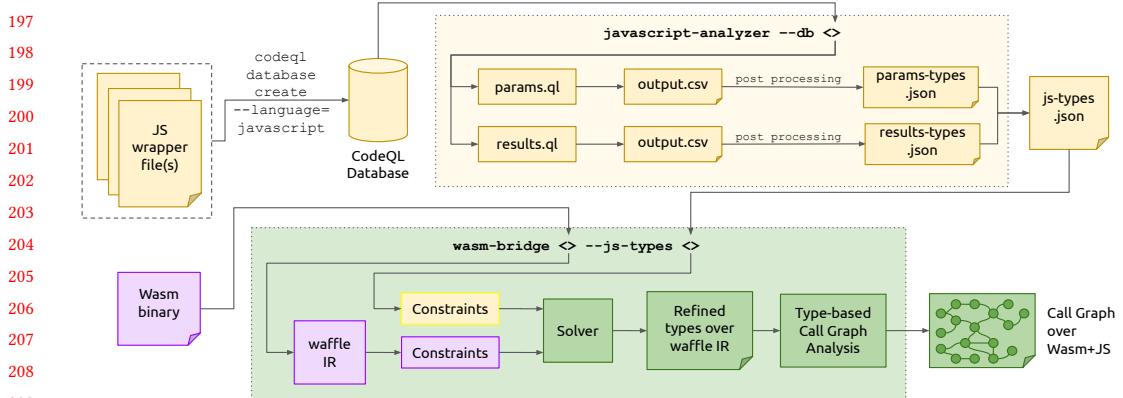


Fig. 2. System Diagram for Multi-Language Static Analysis for WebAssembly and JavaScript.

which function to call. The precision of a call-graph analysis depends on how indirect calls are handled. Current state-of-the-art analysis tools handle indirect calls in one of three ways:

- (1) *Naive Analysis*: The target of an indirect call could be any function in the WebAssembly function table. Industry tools like WASM-OPT and WASM-METADCE[41] take this approach.
- (2) *Type-based Analysis*: The syntax of a `call_indirect` instruction contains a function type annotation (line W11) which is guaranteed to match the function type of the indirect call target. State-of-the-art analysis tools like WASSAIL[36] and WASMA[5] restrict the set of possible targets to be the functions in the table whose type matches the type annotation.
- (3) *Index Analysis*: Since an `i32` value is used to index into the function table, tools like WASIILY [31] and STURDY [22] perform a sophisticated value analysis to determine the set of possible targets.

Unfortunately, for the example in Figure 1, all three strategies yield the same result. The indirect calls at W11 and W16 have the same type annotation: `type0`, a function type that takes two `i32`s as arguments and returns an `i32`. The number and array add and multiply operations all compile down to have the same function type `type0` and the WebAssembly table contains functions for all of these operations. Hence, a static analysis over this binary would resolve all four operations to be potential targets for both indirect calls. With a closed-world (WebAssembly only) analysis, there is no way to distinguish `num_add` from `arr_add` since they both have the same type. However, the JavaScript `num_add` and `arr_add` functions do have different types. We can see that `num_add` takes two numbers and passes them to `num_add` as two `i32`s, whereas `arr_add` takes an `Int32Array` and then calls `arr_add` with two `i32`s that actually represent, respectively, a pointer to the array allocated in WebAssembly memory and the length of that array. If our analysis included an analysis of the JavaScript wrapper, we could differentiate between these two Wasm functions more precisely!

2.3 Refining the WebAssembly Type System With Pointers

WebAssembly does not distinguish between numbers and pointers. The WebAssembly memory is only indexed by `i32` values, or 32-bit integers, while `i32` values are also used as conditionals for if-then-else expressions and as indices into the function table. Pointers to the WebAssembly memory can be distinguished in the JavaScript wrapper from regular numbers by virtue of being passed into WebAssembly through a different pattern, and so, we can refine `i32` to two new types, `ptr` and `num`, both subtypes of `i32`. However, we must propagate these refined types through the WebAssembly code to have any hope of improving the precision at an indirect call site. To this end, we modify the WebAssembly typing rules, change the type system to be a refinement type

246 system [9, 16, 19, 33] and type the WebAssembly memory to be able to track reads and writes to
 247 memory. The changes to the type system are described in Section 3. The changes we make to the
 248 type system are not extensive. We only refine i32’s and do not refine the other base WebAssembly
 249 types, i64, f32 and f64. We also do not incorporate any of the rich JavaScript types for data allocated
 250 in the WebAssembly memory into the type system, leaving that to future work. Our hypothesis is
 251 that in working with even a simple set of refined types for WebAssembly, *we are able to be more*
 252 *precise in our estimation of indirect call targets simply by virtue of richer type information.*

253 Figure 2 illustrates how we discover these richer refinement types for WebAssembly. In order
 254 to type the arguments and returns for calls to WebAssembly exported functions in the JavaScript
 255 wrapper, we write a CodeQL [1] analysis that infers a refined type by checking for specific patterns
 256 of accessing the WebAssembly memory. CodeQL is a framework for writing static analyses with
 257 datalog-like queries. We then pass these types into our tool WASMBRIDGE that encodes them as
 258 constraints. We also generate constraints over the WebAssembly binary under inspection. This
 259 allows us to discover refined types – for instance, if a WebAssembly instruction loads from or
 260 stores to an i32 then that type should be refined to a pointer. Instead of working over the binary
 261 directly, we generate constraints over WAFFLE [8], an intermediate representation for WebAssembly
 262 that uses Static Single Assignment Form [10]. We then solve the constraints, while doing a value
 263 and pointer analysis, to get refinement types over the Wasm binary. We then apply a type-based
 264 call-graph analysis over the refined type system. We explain the details of this process in Section
 265 3, 4 and 5. Note that we only generate and solve constraints over WebAssembly 1.0 and support
 266 analysis over JavaScript wrapper files generated by the Rust compiler.

268 3 A REFINEMENT TYPE SYSTEM FOR WEBASSEMBLY

269 3.1 Background on WebAssembly Instructions and Typing

270 Below we discuss a subset of WebAssembly instructions and their typing. For a description of the
 271 complete instruction set and module typing, please refer to Haas et. al.[20] and the WebAssembly
 272 1.0 specification [40]. Instructions in WebAssembly operate on an implicit value stack by popping
 273 argument values and pushing computed results. As shown in Figure 3, the instruction typing
 274 judgment in WebAssembly has the form $C^w \vdash e^* : \tau_w^m \rightarrow \tau_w^n$ which says that in a context C^w ,
 275 the one-or-more instructions e^* expect a sequence of m values of types τ_w^i (for $i \in \{1, \dots, m\}$) at the
 276 top of the stack and replace these with a sequence of n values of types $\tau_w'^j$ (for $j \in \{1, \dots, n\}$) at the
 277 top of the stack. Note that here we use subscripts and superscripts w for grammars of original
 278 WebAssembly elements, while later we will use r for our refinement-typed WebAssembly.
 279

280 *3.1.1 Values and Arithmetic Instructions.* WebAssembly has four value types: i32, i64, f32, f64. They
 281 represent 32- and 64-bit integers and 32- and 64-bit floating-point numbers. Values in WebAssem-
 282 bly are numeric constants tagged with the appropriate value type, e.g., the 32-bit integer 42 is
 283 represented as i32.const 42. The i32.const 42 instruction pops nothing from the stack and pushes
 284 the i32.const 42 value onto the stack. This is reflected in the typing rule for constants, as shown in
 285 Figure 3. The figure also contains the typing rule for binary operations, which expect two values of
 286 the type they are annotated with on the stack and produce a value for the same type. For example,
 287 i32.add expects two i32’s on the stack and pushes a i32 onto the stack.

288 *3.1.2 Control Constructs and Breaks.* WebAssembly has control constructs such as blocks and
 289 loops and provides structured control flow with break instructions that are annotated with an index:
 290 br i . Here, i is a de-bruijn index [11] out of n labels that are associated with n enclosing control
 291 constructs. The target construct at i determines how the br instruction behaves –if it is a block,
 292 control jumps to the *end* of the block and if it is a loop, control jumps to the *start* of a loop. This is
 293

295 $C^w ::= \{\text{func } tf_w^*, \text{ local } \tau_w^*, \text{ global } \tau_w^*, \text{ table } n^?, \text{ memory } n^?, \text{ label } (\tau_w^*)^*, \text{ return } (\tau_w^*)^?\}$

296
297 **Typing Instructions With the Original WebAssembly Value Types.**

$C^w \vdash e^* : tf_w^*$

$$\tau_w ::= i32 | i64 | f32 | f64$$

$$tf_w ::= \tau_w^* \rightarrow \tau_w^*$$

$$\frac{}{C^w \vdash \tau_w.\text{const } c : \epsilon \rightarrow \tau_w} \text{CONSTANT}$$

$$\frac{}{C^w \vdash \tau_w.\text{binop} : \tau_w \tau_w \rightarrow \tau_w} \text{BINARY OPS}$$

$$\frac{tf_w = \tau_w^m \rightarrow \tau_w^n \quad C^w, \text{label}(\tau_w^n) \vdash e^* : tf_w}{C^w \vdash \text{block } tf_w e^* \text{ end} : tf_w} \text{BLOCK}$$

$$\frac{C_{\text{label}}^w(i) = \tau_w^m}{C^w \vdash \text{br } i : \tau_w^* \tau_w^m \rightarrow \tau_w^*} \text{BREAK}$$

$$\frac{C_{\text{func}}^w(i) = tf_w \quad C_{\text{memory}}^w = n \quad 2^a \leq (|tp| <)^? |\tau_w| \quad (tp_sz)^? = \epsilon \vee \tau_w = \text{im}}{C^w \vdash \text{call } i : tf_w} \text{CALL}$$

$$\frac{}{C^w \vdash \tau_w.\text{load}(tp_sx)^? a o : i32 \rightarrow \tau_w} \text{LOAD}$$

310 Fig. 3. Typing of a subset of WebAssembly Instructions in the Original Type System.

311 evident in the typing of the block and br instructions. Blocks are annotated with a signature tf_w
 312 that expects τ_w^m types on the top of the stack at the start of the block, and expects τ_w^n types on the
 313 top of the stack at the end of a block. The sequence of instructions in a block is type checked under
 314 a label that states that this enclosing block expects τ_w^n values on the stack. A br instruction finds
 315 the label associated with the i^{th} enclosing target control construct and before jumping to the end
 316 or start of the target, ensures it has the right types on the stack. If the target construct is a block, it
 317 would expect τ_w^n on the top of the stack.
 318

319 **3.1.3 Loads and Stores to Linear Memory.** Loads and stores to WebAssembly memory are done
 320 with load and store instructions that are annotated with the type of data loaded from memory and
 321 stored in memory respectively. For example, i64.load loads a i64 value from memory. WebAssembly
 322 memory is only indexed by i32 values and so the i64.load expects an i32 value on the stack and
 323 pushes an i64 value to the stack.

324 **3.1.4 Function Calls.** Functions in WebAssembly are referenced using an immediate index into the
 325 function section of a module. Hence, a call instruction is annotated with a index i , identifying the
 326 function to be called. This index is used to look up the function type of the called function in the
 327 context C^w . The function type specifies the types at the top of the stack before and after the call.
 328

329 3.2 Refinement Type System

330 We present a refinement over the original WebAssembly (version 1.0) type system in Figure 4. Here,
 331 R is a refined store context, analogous to the store context S in the standard WebAssembly type
 332 system. We discuss each change to the type system below:

- 334 (1) *Refinement Types, τ_r :* We introduce refinements to each base WebAssembly value type τ_w .
 335 For example, if the top of the stack is typed to be τ_w in WebAssembly, it is ascribed the
 336 refinement type $\tau_w(n)$ in our type system. A refinement type such as $i32(42)$, in a more
 337 traditional presentation of refinement types, might be written as $\{n : \beta \mid \beta <: i32 \wedge n = 42\}$.
 338 For types $i32(n)$ and $f32(n)$, $n \in \{\mathbb{N}_{32} \cup \top\}$, while for $i64(n)$ and $f64(n)$, $n \in \{\mathbb{N}_{64} \cup \top\}$.
 339 Floating point numbers are represented by their bits.
- 340 (2) *Pointers and Numbers:* Since i32s are used both as pointers into the WebAssembly memory
 341 and numbers, we add two new types, $\text{ptr}(n)$ and $\text{num}(n)$, where ptr and num are subtypes
 342 of $i32$. The types $i32(n)$, $\text{ptr}(n)$, and $\text{num}(n)$ form a lattice, the subtyping relation and
 343

344 $R ::= \{\text{inst } C^*, \text{ tab } n^*, \text{ mem } \Psi^*\}$
 345 $C^r ::= \{\text{func } tf_r^*, \text{ local } \tau_r^*, \text{ global } \tau_r^*, \text{ table } n^?, \text{ memory } n^?, \text{ label } (\tau_r^*, \Psi)^*, \text{ return } (\tau_r^*, \Psi)^?\}$

Typing WebAssembly Instructions with Refinement Types

 $R; C^r \vdash e_r^* : tf_r$

$$\begin{array}{c}
 \begin{array}{l}
 \tau_{r32} ::= i32(n) \mid \text{ptr}(n) \mid \text{num}(n) \mid f32(n) \\
 \tau_r ::= \tau_{r32} \mid i64(n) \mid f64(n) \\
 \tau ::= i32 \mid \text{ptr} \mid \text{num} \mid i64 \mid f32 \mid f64 \\
 \Psi ::= \{(\text{ptr}(n) \mapsto \tau_{r32})^*\} \\
 tf_r ::= \tau_r^*, \Psi \rightarrow \tau_r^*, \Psi'
 \end{array}
 \\[10pt]
 \frac{R; C^r \vdash \tau.\text{const } c : \tau(c)}{R; C^r \vdash \tau.\text{const } c : \tau(c)} \text{ CONSTANT} \quad \frac{\tau_r^3 = \text{binop}(\tau_r^1, \tau_r^2) \quad \tau_r^3 <: \tau(\top)}{R; C^r \vdash \tau.\text{binop} : \tau_r^1 \tau_r^2 \rightarrow \tau_r^3} \text{ BINARY OPS}
 \end{array}$$

$$\frac{\begin{array}{c} tf_r = \tau_r^m, \Psi_{\text{pre}} \rightarrow \tau_r^n, \Psi_{\text{post}} \quad R\Psi = \Psi_{\text{pre}} \\ R; C^r, \text{label}(\tau_r^n, \Psi_{\text{post}}) \vdash e^* : \tau_r^m, \Psi_{\text{pre}} \rightarrow (\tau_r')^n, \Psi'_{\text{post}} \\ (\tau_r' <: \tau_r)^n \quad \Psi'_{\text{post}} <: \Psi_{\text{post}} \end{array}}{R; C^r \vdash \text{block } tf_r \ e^* \ \text{end} : tf_r} \text{ BLOCK} \quad \frac{\begin{array}{c} C_{\text{label}}^r(i) = (\tau_r')^m, \Psi' \\ (\tau_r <: \tau_r')^m \quad R\Psi <: \Psi' \end{array}}{R; C^r \vdash \text{br } i : \tau_r^* \ \tau_r^m \rightarrow \tau_r^*} \text{ BREAK}$$

$$\frac{\begin{array}{c} C_{\text{func}}^r(i) = (\tau_r')^m, \Psi'_{\text{pre}} \rightarrow (\tau_r')^n, \Psi'_{\text{post}} \\ tf_r = \tau_r^m, \Psi_{\text{pre}} \rightarrow \tau_r^n, \Psi_{\text{post}} \\ R\Psi = \Psi_{\text{pre}} \quad (\tau_r <: \tau_r')^m \quad (\tau_r' <: \tau_r)^n \\ \Psi_{\text{pre}} <: \Psi'_{\text{pre}} \quad \Psi'_{\text{post}} <: \Psi_{\text{post}} \end{array}}{R; C^r \vdash \text{call } i \ tf_r : \tau_r^m, \Psi_{\text{pre}} \rightarrow \tau_r^n} \text{ CALL} \quad \frac{\begin{array}{c} C_{\text{memory}}^r = n \\ \tau_r = \text{load_and_extend}(c + o, \tau, (tp_sx)^?, R\Psi) \\ 2^a \leq (|tp| <)^? |\tau_r| \quad (tp_sz)^? = e \vee \tau_r = im \end{array}}{\Psi; C^r \vdash \tau.\text{load}(tp_sx)^? \ a \ o : \text{ptr}(c) \rightarrow \tau_r} \text{ LOAD}$$

Fig. 4. Typing of a subset of WebAssembly Instructions in the Refinement Type System.

368 meet and join operations for which are discussed in Section C. Refinements of the other
 369 WebAssembly value types, $i64(n)$, $f32(n)$, and $f64(n)$, are not refined further since they are
 370 not used to index memory.

- 371 (3) *Memory Typing, Ψ :* Unlike the original WebAssembly type system, we type the flat contiguous array of bytes that serves as the WebAssembly linear memory. Since we hope to propagate types from JavaScript to indirect call sites, and since memory operations are ubiquitous in WebAssembly, we need to be able to recover the (refined) types of values that are stored in memory at a specific address. A memory typing Ψ is a mapping from pointer addresses a , which are 32-bit integers, to refinement types τ_r .
- 372 (4) *Function Types:* Function types tf_w in WebAssembly are used to encode the types of functions, blocks, loops, and calls. They specify the (input and output) base WebAssembly types τ_w for all of these constructs. In our refinement type system, a function type specifies not only the lists of (input and output) refinement types τ_r , but also the shape of the (input and output) memory typing Ψ . We must track how the memory typing changes over the course of execution since WebAssembly programs perform strong updates that can change the type of data stored in memory. For instance, a function that expects a pointer at location a , because it loads from a , might later store an $f32$ at location a .
- 373 (5) *Label Typing:* In WebAssembly, labels of control constructs contain the shape of the stack expected when there is a jump to the control construct. In our refinement type system, 374 we must specify not only the expected refined types τ_r on the stack but also the expected 375 memory typing Ψ for the control construct in the label.

393 Next, we describe the refinement typing rules for a subset of WebAssembly instructions (shown
 394 in Figure 4). Refinement typing rules for the entire instruction set can be found in Appendix A.

395

396 **3.2.1 Arithmetic Instructions.** We restrict the refinement typing of arithmetic instructions to only
 397 allow certain combinations of i32, ptr and num. For instance, in the BINARY OPS rule in Figure 4,
 398 the result type is computed using the function binop which takes two input refinement types and
 399 computes a result refinement type, but only considers certain pairs of input types valid for each
 400 specific binary instruction. The instruction will fail to type check if a invalid pair of input types is
 401 provided for a given binary operation. We discuss a few specific binary operations below.

402

- 403 • **Comparison:** Comparison of i32s and all its subtypes result in a num. All combinations of
 404 i32, ptr and num can be compared.
- 405 • **Addition:** WebAssembly allows addition of two i32s, either of which, or the result of which,
 406 can be used to index into the WebAssembly memory. We do not allow addition of two
 407 pointers, but allow addition of every other combination of i32, ptr and num. Moreover,
 408 addition of a ptr and a num yields a ptr.
- 409 • **Subtraction:** We do not allow subtraction of ptr from a num, but we do allow all other i32,
 410 ptr and num combinations for the subtraction operation. Subtracting a ptr from a ptr yields
 411 a num (an offset).

412

413 **3.2.2 Control Constructs and Breaks.** The BLOCK rule is annotated with a refined function type
 414 tf_f which specifies the stack shape and memory typing expected before the block instruction and
 415 at the end of the block instruction. Since the end of a block is the meet of several breaks out of
 416 the block, we ensure that the shape of the stack and memory after type checking the instructions
 417 in the block should be subtypes of the types expected on the top of the stack and of the expected
 418 memory typing. This is necessary since it is valid in WebAssembly for one control path to the end
 419 of the block to return a ptr and another to return a i32, since they would both be integers in the
 420 WebAssembly type system. br instructions are typed similarly.

421

422 **3.2.3 Loads from Linear Memory.** While in WebAssembly, i32s are used to index into memory, our
 423 refinement type system requires that only values of ptr type be used to index into memory. We
 424 modify the typing rule for the LOAD instruction to expect a $\text{ptr}(n)$ on the stack, as shown in Figure
 425 4. The load instruction is optionally annotated with a packed type and size, tp_{sx} , which is used to
 426 pack and sign extend the data stored in memory. Additionally, our memory typing Ψ only stores
 427 types of 32-bit size, $\tau_{r_{32}}$. If the load instruction expects a type of size 64 bits, we load the sequence of
 428 bytes from address $n + o$, where o is the offset provided to the load instruction, and address $n + o + 4$
 429 as a 64-bit value as directed by the τ annotation on the load instruction which specifies the type of
 430 data we want to load. All this is done by the `load_and_extend` function, in Appendix A.

431

432 **3.2.4 Function Calls.** Typing a call instruction is more complex than in WebAssembly since our
 433 refinement type system has subtyping and since it keeps track of the expected memory typing at
 434 various points in the program. We change the WebAssembly `call i` instruction to `call i tf_r`, where
 435 tf_r specifies the stack shape and memory typing expected before and after the call. When typing
 436 the call instruction, we need to ensure that the stack shape and memory typing before the call are
 437 subtypes of the stack shape and memory typing expected by the function being called, and that the
 438 stack shape and memory typing when the function returns are subtypes of the stack shape and
 439 memory typing expected after the call by the callee.

440

442	Refinement Type System	
443		
444	<i>Symbolic Number</i>	
445	<i>Symbolic Location</i>	
446	<i>Abstract Refinement Type</i>	
447	$\hat{\tau}_r$	$::= \text{i32}(n) \mid \text{ptr}(l, n) \mid \text{num}(n) \mid \text{i64}(n) \mid \text{f32}(n) \mid \text{f64}(n)$
448	τ	$::= \text{i32} \mid \text{ptr} \mid \text{num} \mid \text{i64} \mid \text{f32} \mid \text{f64}$
449	<i>Original Value Types</i>	
450	τ_w	$::= \text{i32} \mid \text{i64} \mid \text{f32} \mid \text{f64}$
451	<i>Symbolic Function Type</i>	
452	tf_α	$::= \alpha^* \rightarrow \alpha^*$
453	<i>Symbolic Global Type</i>	
454	tg_α	$::= \text{mut? } \alpha^*$
455	<i>Abstract Memory Typing</i>	
456	ς	
457	<i>Symbolic Memory Typing</i>	
458	Σ	$::= \cdot \mid \Sigma, (l, n) \mapsto \hat{\tau}_r$
459	Constraint System	
460		
461	<i>Operations</i>	
462	op	$::= \text{unop}_{iN} \mid \text{unop}_{fN} \mid \text{ctop} \mid \text{binop}_{iN} \mid \text{binop}_{fN} \mid \text{testop}_{iN} \mid \text{relop}_{iN} \mid \text{relop}_{fN}$
463	<i>Type Constraint</i>	
464	t	$::= \alpha \doteq \hat{\tau}_r \mid \alpha <: \tau \mid \alpha \doteq \alpha \mid l \neq l \mid \alpha \doteq \sqcup \alpha^+ \mid \alpha \doteq \varsigma[\alpha]_{(tp_sx^?, a, o)} \mid \alpha \doteq op((\alpha^?)^+)$
465	<i>Memory Constraint</i>	
466	m	$::= \varsigma \doteq \Sigma \mid \varsigma \doteq \varsigma \mid \varsigma \doteq \sqcup \varsigma^+ \mid \varsigma \doteq \varsigma[\alpha \mapsto \alpha]_{(tp^?, a, o)}$
467	<i>Constraint Set</i>	
468	S	$::= \cdot \mid S, t \mid S, m$

Fig. 5. Syntax for Symbolic Refinement Types and Constraints Generated Over WebAssembly.

3.3 Discovering Refinement Types using Type and Memory Constraints

In this section, we explain how we discover refinement types for WebAssembly code. WebAssembly's type validation algorithm, described in the specification [40], performs a single forward pass over the instructions in a function body to validate that WebAssembly function. Unfortunately, a forward-only analysis isn't possible when we wish to discover refinement types. Instead, we first generate constraints over a WebAssembly function in order to later discover refinement types for each function, instruction, and variable via constraint solving. Accumulating constraints is necessary because it is often not immediately obvious if an i32 value is a pointer or a number.

Another point of interest is that pointers in WebAssembly don't often have a concrete memory address that can be statically known. Our refinement type system expects the pointer type to have a concrete address a as its value — written $\text{ptr}(a)$ — and for memory typing Ψ to be a mapping from concrete addresses a to values. Note that if the address is unknown, we would have to resort to \top as the address value, which reduces precision. Thus, when discovering refinement types, we introduce *symbolic* pointer types $\text{ptr}(l, n)$ that have a symbolic base address l and a symbolic offset n . With symbolic pointers in hand, we define a grammar for symbolic refinement types $\hat{\tau}_r$, given in Figure 5. We also introduce symbolic memory typing to be a map from symbolic pointers $\text{ptr}(l, n)$ to symbolic refinement types $\hat{\tau}_r$. We generate constraints over WebAssembly functions and solve them to get *symbolic* refinement types $\hat{\tau}_r$ for every stack slot in the value stack, and local and global variables, and *symbolic* memory typing Σ . Our call-graph analysis is over the symbolic refinement type system. We will later use the symbolic refinement types we discover for WebAssembly code e^* to guide the construction and typing of a concrete-refinement-typed WebAssembly program e_r^* , for which we prove type safety in Section 3.4.

The symbolic refinement types and memory typing, as well as the constraints that are generated over WebAssembly instructions can be found in Figure 5. We associate every stack slot in the value stack and local and global variables, with a unique abstract refinement type variable α and every memory state with a unique abstract memory variable ς . We generate constraints over

491 $C^\alpha ::= \{\text{func } tf_w^*, \text{ local } \alpha^*, \text{ global } tg_\alpha^*, \text{ table } n?, \text{ memory } n?, \text{ label } (\alpha^*, \varsigma)^*, \text{ return } (\alpha^*, \varsigma)^?\}$

492 **Constraint Generation for Instructions**

$S; \varsigma; C^\alpha \vdash e : \alpha^* \rightarrow \alpha^*; S'; \varsigma'$

493

$$\frac{\alpha \text{ fresh} \quad S' = S :: [\alpha \doteq \tau_w(c)]}{S; \varsigma; C^\alpha \vdash \tau_w.\text{const } c : \epsilon \rightarrow \alpha; S'; \varsigma} \text{CONSTANT}$$

494

$$\frac{\alpha_1 \alpha_2 \in \text{dom}(S) \quad \alpha_3 \text{ fresh}}{S' = S :: [\alpha_1 \doteq \text{binop}(_, \alpha_2, \alpha_3) \wedge \alpha_2 \doteq \text{binop}(\alpha_1, _, \alpha_3) \wedge \alpha_3 \doteq \text{binop}(\alpha_1, \alpha_2, _)])} \text{BINARY OPS}$$

500

$$\frac{\begin{array}{c} \alpha^n \in \text{dom}(S) \quad \alpha^m, \varsigma' \text{ fresh} \\ S; \varsigma; C, \text{label}(\alpha^m, \varsigma') \vdash e^* : \alpha^n \rightarrow (\alpha')^m; S'; \varsigma'' \\ S'' = S' :: [(\alpha \doteq \alpha')^m \wedge \varsigma' \doteq \varsigma''] \end{array}}{S; \varsigma; C^\alpha \vdash \text{block } tf_w e^* \text{ end} : \alpha^n \rightarrow \alpha^m; S''; \varsigma'} \text{BLOCK}$$

505

$$\frac{\alpha^* \alpha^n \in \text{dom}(S) \quad C_{\text{label}}^\alpha(i) = \alpha'_n, \varsigma'}{S' = S :: [\varsigma \doteq \sqcup \varsigma' \wedge (\alpha \doteq \sqcup \alpha \alpha')^n]} \text{BREAK}$$

506

$$\frac{\begin{array}{c} C_{\text{func}}^\alpha = \tau_w^m \rightarrow \tau_w^n \quad \alpha^n \in \text{dom}(S) \\ S' = S :: [(\alpha <: \tau_w)^m \wedge (\alpha' <: \tau_w)^n] \end{array}}{S; \varsigma; C^\alpha \vdash \text{call} : \alpha^m \rightarrow (\alpha')^n; S'; \varsigma} \text{CALL}$$

$$\frac{\alpha_1 \in \text{dom}(S)}{S' = S :: [\alpha_1 <: \text{ptr} \wedge \alpha_2 \doteq \varsigma[\alpha_1](tp_sx?; a, o) \wedge \alpha_2 <: \tau_w]} \text{LOAD}$$

509

510 Fig. 6. Constraint-Generation Rules for a Subset of WebAssembly Instructions.

511

512 WAFFLE [8], an SSA IR over WebAssembly. For each instruction, we generate *type constraints* over α s and *memory constraints* over ς s, the syntax of which can be seen in Figure 5. A subset of the constraint-generation rules for WebAssembly instructions are described in Figure 6, with all the rules in Appendix D. We discuss the interesting constraint-generation cases below.

517

518 **3.3.1 Arithmetic Instructions.** Let us consider the example of i32.add, which expects two values
519 on the stack, of types α_1 and α_2 and produces a single value of type α_3 , i.e., $\alpha_3 = \alpha_1 + \alpha_2$. We
520 know that the only disallowed case for this operation is the addition of two pointers. This has
521 interesting implications for constraint solving, since if $\alpha_1 = \text{ptr}(l, n_1)$ and $\alpha_2 = \text{i32}(n_2)$, α_2 can be
522 refined to be num(n_2), since the alternative is not permitted. In fact, the types of each argument
523 and result can affect each other. If $\alpha_1 = \text{num}(n_1)$, $\alpha_2 = \text{i32}(n_2)$ and $\alpha_3 = \text{ptr}(l, n_3)$, α_2 is refined to
524 be ptr. Meanwhile, if $\alpha_3 = \text{num}(n_3)$, α_2 is refined to num. We leave discovering the appropriate
525 refinements to constraint solving and at the time of constraint generation, generate constraints
526 to tie each α to each other α , with the relevant arithmetic instruction. For i32.add, we generate
527 the constraints, $\alpha_1 \doteq \text{i32.add}(_, \alpha_2, \alpha_3)$, $\alpha_2 \doteq \text{i32.add}(\alpha_1, _, \alpha_3)$ and $\alpha_3 \doteq \text{i32.add}(\alpha_1, \alpha_2, _)$. The
528 underscore in the constraint holds the place of the α current being refined. A similar strategy is
529 followed for all arithmetic instructions.

530

531 **3.3.2 Memory Instructions.** The i64.store instruction expects a pointer type α_{ptr} on the stack
532 and some data type α_{data} . α_{ptr} is constrained to be a pointer by generating a subtyping constraint
533 $\alpha_{\text{ptr}} <: \text{ptr}$, while the data is constrained to be a subtype of the type expected by the store instruction,
534 $\alpha_{\text{data}} <: \text{i64}$. The abstract memory typing represented by ς , an input to the rule, is updated to record
535 a mapping from $\alpha_{\text{ptr}} \mapsto \alpha_{\text{data}}$, and is copied into a new ς' , which is returned by the rule. The
536 constraint on the abstract memory typing is $\varsigma' = \varsigma[\alpha_{\text{ptr}} \mapsto \alpha_{\text{data}}]$. The i64.load instruction is
537 similarly constrained to expect a pointer on the stack and a subtype of i64 as the result. Here, the
538 load constraint is represented as $\alpha_{\text{data}} \doteq \varsigma[\alpha_{\text{ptr}}]$. Similar constraints are generated for all $\tau_w.\text{load}$
539 and $\tau_w.\text{store}$ instructions.

540 3.3.3 *Blocks, Loops.* Let us consider the sequence of instructions, block $tf_w e^*$ end. Let the block
 541 expect n values on the stack and return m values. $m \alpha_1$'s are freshly generated and added to the
 542 context with a label. This is done so that break instructions out of blocks know the number of
 543 values required on the output stack of the target block. We generate constraints for the body of the
 544 block, with this new context, which gives us $m \alpha_2$ s. The $m \alpha_1$ and α_2 s are now equated with equality
 545 constraints, and $m \alpha_1$ s are returned by the instruction on the result stack. A similar scenario arises
 546 for loop instructions. For loops, the n input α s are added to the context instead of the m result α s
 547 that are added to the context for blocks. This is because break instructions in loops restart the loop
 548 and have to know the number of values required on the input stack of the loop.
 549

550 3.3.4 *Break Instructions.* br instructions are used to break out of blocks and restart loops. In the
 551 case of blocks, the values on the stack at the br instruction are returned by the target block. Since
 552 our analysis is flow-insensitive, we join all possible result stacks for a block. The context carries the
 553 expected value stack at the target block, so we generate join constraints for the α s on the stack and
 554 the α s expected by the target block. In the case of loops, the values on the stack at the br instruction
 555 are used to restart the loop. Constraint generation is the same regardless of whether the target of
 556 a br instruction is a loop or a block. Note that generating join constraints in this fashion embeds
 557 the looping structure of the program into the constraints. For example, consider the following
 558 sequence of instructions: loop $tf_w e^*$ end. If $tf_w = i32 \rightarrow ()$, we would generate constraints for
 559 the loop inputs as $\alpha_0 :: i32$. We would then generate constraints for the body of the loop with
 560 label(α_0) added to the context. If we came across the instruction br 0, in the loop body, with α_1 on
 561 the stack, we would generate the constraint $\alpha_0 \doteq \sqcup \alpha_0 \alpha_1$.
 562

563 3.4 Refining WebAssembly and Type Safety for the Refinement Type System

564 In this section, we show two results. First, we explain how, given a WebAssembly program, potentially
 565 in a JavaScript context, we can discover symbolic refinement types and use them to arrive at a concrete-refinement-typed WebAssembly program. Second, we prove type safety for our refinement
 566 type system using progress and preservation.
 567

568 From WebAssembly to Refinement-Typed WebAssembly. We've seen how to generate constraints
 569 over a given WebAssembly function to discover *symbolic* refinement types $\hat{\tau}_r$. However, our typing
 570 rules are over concrete refinement types τ_r and to show that a WebAssembly program is refinement-typed we need to use the symbolic refinement types and symbolic memory typing we
 571 have discovered to guide the annotation of the WebAssembly program with refinement types so
 572 that type-checks using the refinement type system. In order to obtain the concrete refinement types
 573 for a module, we first generate constraints for a given function. During constraint generation, we
 574 map each expression at every point of the program to the α s and ς s it generates on the stack. We
 575 elide this in Figure 6 for conciseness. Section 5 describes how we solve constraints over α s and ς s
 576 to get a least solution ρ that maps α s to symbolic refinement types $\hat{\tau}_r$ and ς s to symbolic memory
 577 typing Σ . However, these refinement types are still symbolic. After constraint solving, we posit
 578 that there exists a mapping η from symbolic addresses to concrete addresses for every Σ , where,
 579 for a Σ_{pre} before and Σ_{post} after an instruction, $\eta_{\text{post}} \supseteq \eta_{\text{pre}}$. Hence, for every α and ς , we recover a
 580 τ_r and Ψ . Since Ψ is only typed with 32-bit types, we split each 64-bit sized type $i64(n)$ or $f64(n)$,
 581 into two $i32$ values from the high and low 32 bits of n . We now make several transformations to
 582 the WebAssembly function, in order to type check it:
 583

- 584 • Instructions tagged with a WebAssembly type annotation τ_w are transformed to have a refined type tag τ . E.g. $i32.\text{const} 42 \rightsquigarrow \text{num}.\text{const} 42$, where the latter expects a num on the stack rather than a $i32$.

- Instructions annotated with WebAssembly function type annotations tf_w are instead annotated with concrete refinement function types tf_r that include Ψ_{pre} and Ψ_{post} memory typing before and after the instruction.
- Function types are also transformed to have concrete refinement function types tf_r .
- The call instruction is typed to have a tf_r annotation (shown in Figure 4) to denote the shape of the stack before and after the call.

Type Safety for the Refinement Type System. The operational semantics of this transformed WebAssembly program e_r^* remains largely unchanged from the original WebAssembly operational semantics, but there are a few changes. For example, for the binary operation rule, the type annotation on the $\tau.w.\text{binop}$ instruction no longer tells us the types expected on the WebAssembly stack. Instead of, $(\tau_w.\text{const } c_1)(\tau_w.\text{const } c_2) \tau_w.\text{binop} \hookrightarrow \tau_w.\text{const}(\text{binop}(c_1, c_2))$, in the refinement type system, the operational semantics of binop steps as follows, $(\tau_1.\text{const } c_1) (\tau_2.\text{const } c_2) \tau_3.\text{binop} \hookrightarrow \text{binop}((\tau_1.\text{const } c_1), (\tau_2.\text{const } c_2))$.

We now prove type safety as the standard progress and preservation theorems.

THEOREM 1. Progress: If $R; C^r \vdash_i e_r^* : \tau_r^1 \cdots \tau_r^m \rightarrow \tau_r^* \wedge (v : \tau_r)^m \wedge \vdash_i s : R$ then $e_r^* = v^*$ or $e_r^* = \text{trap}$ or $s; v_r^*; e_r^* \hookrightarrow s'; v'^*; e_r'^*$

THEOREM 2. Preservation: If $R; C^r \vdash_i e_r^* : \tau_r^1 \cdots \tau_r^m \rightarrow \tau_r^* \wedge \vdash_i (v : \tau_r)^m \wedge \vdash_i s : R \wedge s, v^*, e_r^* \hookrightarrow s', v'^*, e_r'^*$ then $\exists R'$ such that $\text{dom}(R_\Psi) \subseteq \text{dom}(R'_\Psi) \wedge \vdash_i s' : R' \wedge R'; C^r \vdash_i e_r'^* : \tau_r^*$

Proofs of the above theorems are very similar to the WebAssembly progress and preservation proofs except for minor changes to account for the refined types. The main interesting aspect of our refinement type system is that, unlike WebAssembly, we also type the linear memory, keeping track of the memory typing Ψ in the context R . We present the proof of progress and preservation for the load instruction in Appendix B.

4 GENERATING CONSTRAINTS OVER THE JAVASCRIPT WRAPPER

We've seen how to recover refinement types over WebAssembly, but we can also type the parameters passed into WebAssembly exported function calls and the results of these calls to have refinement types. We can analyze the glue code before and after a call to an exported WebAssembly function to type exported functions. Specifically, we want to type arguments as pointers when a JavaScript object is *lowered* to WebAssembly by fetching a pointer from malloc and writing the object's contents to memory at that pointer. We also want to be able to type the data being put into the WebAssembly memory. Similarly, we want to type results as pointers when a WebAssembly value is *lifted* to JavaScript by reading data out of WebAssembly memory, into a JavaScript object. We also want to type the data that is expected to be at that WebAssembly memory location.

To this end, we write a CodeQL analysis to infer refinement types for exported WebAssembly function calls. CodeQL [1] is a declarative semantic code analysis framework maintained by GitHub, in which analysis writers can write datalog-like queries for languages like JavaScript, Python, etc. It is particularly helpful in analyzing code to check for specific patterns. The CodeQL analysis consists of two parts: an analysis to identify calls to the exported `malloc` function and accesses to WebAssembly memory, and a type inference engine that infers a type for a given JavaScript expression. The former is a series of rules that check for a sequence of patterns through data-flow and the latter is a type-inference engine. We rely on type annotations provided by the wrapper file, annotated TypeScript library functions and typing through dataflow expressions to infer a type. For example, to infer the type for a if-then-else expression, we type the then and else branches. We also infer the type of an expression, by trying to infer the type of its parent or, more generally,

Application	Size (KB)	Functions	Exported Functions	Imported Functions	Direct Calls	Indirect Calls	Functions in Table
webcam	28.61	125	6	33	270	79	58
leaflet-rs	20.16	110	8	45	228	39	39
wasm-rsa	369.00	785	25	5	6058	107	237
blake3	34.41	81	14	1	206	26	21
epqueue	23.26	123	17	9	471	25	29
serde	13.61	44	0	1	67	25	31
snowpack-ts	35.11	107	4	37	236	55	35
warp	162.09	283	10	14	1856	112	141
wasm-pack	13.78	43	1	6	69	19	17
source-map	48.53	68	10	1	395	1	18

Table 1. Description of Subject Applications

its predecessor (an expression whose data flows into the current expression) or successor (an expression into which the current expression’s data flows).

The inferred types for each argument and result of exported WebAssembly functions are passed into WASMBRIDGE which generates constraints from them, as shown in Figure 2. Generating constraints from the types is fairly straightforward. If an argument or result is typed to be a pointer with that memory block containing a i64, we generate the subtype constraints $\alpha_1 \doteq \text{ptr}(l_1, 0)$, $\alpha_2 <: \text{i64}$ and the memory constraint $\varsigma_1 \doteq \varsigma_0[\alpha_1 \mapsto \alpha_2]$. Note that since we do not analyze the JavaScript client, we do not have any data to associate with these types. If a function has two arguments that have a pointer type and we assume that a call to malloc gives us fresh, disjoint labels, we can generate a constraint that encodes that the two pointers’ symbolic base addresses are not equal to each other. For example, if func \$foo (α_1, α_2) where $\alpha_1 \doteq \text{ptr}(l_1, 0)$ and $\alpha_2 \doteq \text{ptr}(l_2, 0)$, we generate the constraint $l_1 \neq l_2$ since our analysis over the JavaScript wrapper tells us that α_1 and α_2 came from two separate malloc calls. We solve these constraints along with the constraints generated over the associated WebAssembly binary to obtain refinement types for each WebAssembly function, stack slot and local variable.

5 CONSTRAINT SOLVING

We solve a set of constraints generated over a WebAssembly function and its associated JavaScript glue code using a standard fixpoint algorithm, described in detail in Appendix F. At the end of constraint generation, we have a set S of constraints over a function body. We create a mapping Constraint , that maps each α in S to the type constraints on α and similarly, maps each ς in S to the memory constraints on ς . We resolve the equality constraints for α ’s and ς ’s when creating this map. We also record dependencies between α ’s and ς ’s using the influence vector infl . This is used to recompute the constraints of a certain α or ς when its dependencies have changed. We pass the Constraint and infl maps as inputs to constraint solving and compute a least solution for every α and ς in S . The constraints on α ’s are solved to get a refinement type, τ_r , and the constraints on ς ’s are solved to get a memory state, Σ . We obtain the least solution using a standard worklist algorithm. If there exists a $\alpha \doteq \tau_r$ in the constraints for α , we initialize the solution for the α to be τ_r . Otherwise, we initialize the solution to be the \top type. We initialize the solution of all ς ’s to be an empty memory state. We initialize the worklist with all α ’s and ς ’s in the domain of the Constraint map, and continue solving till the worklist is empty. If the solution of an α or ς is found to be different from its solution in the previous iteration, we add all α ’s and ς ’s that depend on it into the worklist. We briefly discuss solving type and memory constraints below.

		WebAssembly			Refined Type System						WebAssembly + JavaScript					
		E	U	M	E	U	M	% E Red.	% Call Sites w/ Red.	Avg Red.	E	U	M	% E Red.	% Call Sites w/ Red.	Avg Red.
Applications	webcam	526	7	4	499	7	5	9.28	50.00	16.64	499	7	5	9.28	50.00	16.64
	leaflet-rs	264	9	3	229	11	6	17.78	62.86	36.74	229	11	6	17.78	62.86	36.74
	wasm-rsa	3896	8	2	3696	9	2	6.15	35.35	12.41	3696	9	2	6.15	35.35	12.41
	blake3	202	5	13	195	16	14	5.03	19.23	13.00	195	16	14	5.03	19.23	13.00
	epqueue	287	14	3	284	14	5	1.51	15.00	9.50	284	14	5	1.51	15.00	9.50
	serde	96	43	3	90	43	5	12.50	45.00	23.12	90	43	5	12.50	45.00	23.12
	snowpack-ts	352	45	4	338	45	21	7.37	51.92	30.88	338	45	21	7.37	51.92	30.88
	warp	1832	9	3	1707	11	3	8.73	34.86	13.04	1707	11	3	8.73	34.86	13.04
	wasm-pack	88	35	9	83	35	9	11.11	42.11	14.04	83	35	9	11.11	42.11	14.04
	source-map	126	7	0	126	7	0	0.00	0.00	0.00	126	7	0	0.00	0.00	0.00
	Averages	766.9	18.2	4.4	724.7	19.8	7	7.946	35.633	16.937	724.7	19.8	7	7.946	35.633	16.937

Table 2. Comparing callgraphs generated over the original WebAssembly type system, the refinement type system, and the combined WebAssembly and JavaScript analysis. $|E|$ denotes the number of edges in the callgraph, $|M|$ denotes the number of monomorphic calls, and $|U|$ denotes the number of unreachable nodes. Avg Red. is average reduction at call sites.

5.1 Solving Type Constraints

To solve the type constraints for a given α we iterate over each type constraint and do a meet operation between the currently computed type for the α and the type computed with a specific constraint. The solution of most type constraints is standard: we simply evaluate the arithmetic operations or perform joins as dictated by the join constraints. However, the memory load constraint is more complicated in the face of possible aliasing between pointers. Let us suppose that we are trying to load from a memory state, Σ , at $\text{ptr}(l, n)$. If Σ contains $\text{ptr}(\top, \top)$, $\text{ptr}(l, \top)$ or $\text{ptr}(l', n')$ where there does not exist a $l \neq l'$ constraint, we make the conservative assumption that all these pointers might alias with $\text{ptr}(l, n)$ and join the data pointed to by these potential aliased pointers with τ_r where $\Sigma[(l, n) \mapsto \tau_r]$.

5.2 Solving Memory Constraints

Solving the memory constraints for a given ζ is similar to solving type constraints. Here too, when evaluating a memory store constraint $\Sigma[\text{ptr}(l, n) \mapsto \tau_r]$, we have to be conservative in the face of aliasing and join the data at a potential alias site, i.e., $\text{ptr}(\top, \top)$, $\text{ptr}(l, \top)$ or $\text{ptr}(l', n')$, with the type currently being stored in the memory state, τ_r .

6 EVALUATION ON REAL-WORLD SUBJECT APPLICATIONS

We evaluate a call graph analysis over the refinement type system on ten real-world WebAssembly-/JavaScript applications that have been compiled from Rust. Our approach only supports generating constraints over WebAssembly 1.0, the core WebAssembly language introduced in 2017, and only supports an analysis of JavaScript wrappers generated by the Rust compiler. We discuss extending our analysis to include other compilers and extensions to WebAssembly in Section 8. We pick subject applications from two datasets [28] [38] and identify 29 NPM packages that have been compiled from Rust. Of these, we find 10 libraries which satisfy our criteria of only using WebAssembly 1.0. Table 1 shows the characteristics of the subject applications, including the size of the binary and the number of direct and indirect call sites in the binary. We run WASMBRIDGE over the chosen subject applications in order to evaluate the precision of the refinement types, the analysis over JavaScript and how WASMBRIDGE compares to other state-of-the-art call graph analysis tools.

	WebAssembly Type System	Refined	Wasm+JS Analysis			
		Type System	CodeQL DB Creation	CodeQL Query Execution	Constraint Solving	Total Time (s)
Applications	Time (s)	Time (s)				
webcam	0.005	0.053	4.80	8.33	0.06	13.20
leaflet-rs	0.005	0.024	5.36	7.81	0.03	13.19
wasm-rsa	0.032	1.330	5.14	7.78	1.44	14.36
blake3	0.005	0.045	4.93	7.68	0.05	12.66
epqueue	0.007	0.042	4.29	7.82	0.08	12.18
serde	0.004	0.027	4.82	8.49	0.03	13.34
snowpack-ts	0.007	0.052	5.58	7.38	0.06	13.01
warp	0.023	1.184	4.54	7.97	1.32	13.83
wasm-pack	0.004	0.025	4.46	6.90	0.03	11.39
source-map	0.010	0.189	5.04	7.28	0.20	12.52

Table 3. Runtime performance comparison (in seconds) for analysis over the original WebAssembly type system, the refinement type system, and the combined Wasm+JS analysis for each subject application.

RQ1: How much more precise is the call graph over the refinement type system when compared to the original WebAssembly type system?

In order to determine the increase in precision afforded by the refinement type system, we compare the same type-based call graph analysis on the refinement type system and the original WebAssembly type system. A type-based call graph analysis restricts the set of targets at an indirect call sites with the type annotation at the `call_indirect` instruction as described in Section 2.2. The results of this evaluation are shown in the first two rows of Table 2. The call graph over the refinement type system is more precise than the call graph over the original WebAssembly type system:

- (1) *Edge Reduction*: We reduce the total number of call graph edges by up to 17.50% (`leaflet-rs`) compared to the baseline WebAssembly analysis. This reduction occurs at indirect call sites where precision matters most; we observe a reduction at 62.9% of indirect call sites.
- (2) *Call Site Level Improvements*: The precision gains are more apparent at indirect call sites. For all indirect call sites, we achieve an average edge reduction ranging from 9.5% to 36.74% per call site and an average reduction of 16.9%.
- (3) *Monomorphic Call site Discovery*: Our refinement type system enables us to identify several additional monomorphic call sites—indirect calls that resolve to exactly one target. In `snowpack-ts`, we discover 17 additional monomorphic call sites (from 4 to 21 call sites); `leaflet-rs` has 3 additional monomorphic call sites (from 3 to 6); `webcam-stream` has 1 additional monomorphic call sites (from 4 to 5); `epqueue` has 2 additional monomorphic call sites (from 3 to 5); `serde-bindgen` has 2 additional monomorphic call sites (from 3 to 5) and `blake3` shows 1 additional monomorphic call site (from 13 to 14). These discoveries enable subsequent compiler optimizations and improve the precision of downstream analyses.
- (4) *Unreachable Function Detection*: The refinement type information also improves our ability to identify unreachable functions. `Blake3` shows the most dramatic improvement, with unreachable nodes increasing from 5 to 16—indicating that there are functions in the WebAssembly binary that are never called in practice.

Notably, `source-map` shows no improvement in our analysis because it is a highly optimized binary with only one indirect call. This result is expected in the case of call graph analysis—applications with simple control flow gain little from the refinement type analysis. However, the richer types might improve precision in other analyses.

	WASMBRIDGE			WASM-OPT				WASMA				WASSAIL			
Applications	N	E	U	N	E	U	Edge Diff	N	E	U	Edge Diff	N	E	U	Edge Diff
webcam	125	499	7	125	1903	0	73.78%	125	526	7	5.13%	125	526	7	5.13%
leaflet-rs	110	229	11	110	845	0	72.90%	110	264	9	13.26%	110	264	9	13.26%
wasm-rsa	785	3696	9	785	9501	0	61.10%	785	3896	8	5.13%	785	3896	8	5.13%
blake3	81	195	16	81	269	0	27.51%	78	202	2	3.47%	81	202	5	3.47%
epqueue	123	284	14	123	602	0	52.82%	119	287	10	1.05%	123	287	14	1.05%
serde	44	90	43	44	405	43	77.78%	-	-	-	-	44	96	43	6.25%
snowpack-ts	107	338	45	107	846	0	60.05%	107	352	45	3.98%	107	352	45	3.98%
warp	283	1707	11	283	5537	0	69.15%	280	1832	6	6.82%	283	1832	9	6.82%
wasm-pack	43	83	35	43	153	35	45.75%	43	88	35	5.68%	43	88	35	5.68%
source-map	68	126	7	68	133	0	5.26%	62	126	1	0.00%	68	126	7	0.00%
Average	176.90	724.70	19.80	176.90	2019.40	7.80	54.61%	189.89	841.44	13.67	4.94%	176.90	766.90	18.2	5.08%

Table 4. Call Graph Comparison across different WebAssembly analysis tools. Edge difference shows the % change in callgraph edges compared to our WASMBRIDGE analysis. Negative values indicate the tool produces fewer edges than WASMBRIDGE. |N| denotes the number of nodes (functions) in the callgraph, |E| denotes the number of edges (calls), and |U| denotes the number of unreachable nodes in the callgraph.

RQ2: How much more precise is the call graph over WebAssembly and JavaScript when compared to WebAssembly alone?

We compare call graphs generated over the refinement type system with constraints generated over WebAssembly alone versus constraints generated over WebAssembly and its associated JavaScript wrapper in the second and third row of Table 2. We find no change in the precision at indirect call sites and no reduction in the number of edges. This is not too surprising since the refinements to the type system are not discovered through the JavaScript wrapper alone. WebAssembly instructions also provide these refinements, such as the load instruction which expects a ptr on the stack rather than just an i32. The one constraint that is unique to the JavaScript wrapper is one that specifies the disjointedness of symbolic base addresses of pointers that have come from different calls to malloc ($l_1 \neq l_2$). It seems as though the imprecision of pointer analysis leads to this constraint not improving precision at call sites.

Note that analyzing the wrapper cannot tell us the JavaScript values that are being passed to WebAssembly exported functions. We would need a sophisticated JavaScript client analysis that performs a value and call graph analysis over the entire JavaScript application and all its dependencies, in order to see more precision improvements. No such tool exists off-the-shelf and we leave it to future work. However, the CodeQL analysis of the JavaScript wrapper makes our analysis extensible to such a client analysis. Other state-of-the-art tools cannot support this.

RQ3: What is the cost of the analysis over WebAssembly and JavaScript?

Table 3 shows the runtime overhead of each component of our multi-language analysis and how it compares to the analysis over constraints generated from WebAssembly alone versus the call graph analysis over the original WebAssembly type system. The latter is our baseline. The call graph analysis over the refinement type system over WebAssembly alone shows a maximum of a 1.33 second increase over the baseline. The multi-language analysis, however, is dominated by the time it takes to create a CodeQL database and run the analysis (or query) over the database. Even so, the total time of the multi-language analysis does not exceed 14.36s in the worst case and constraint generation and solving is relatively inexpensive.

Metric		Baseline	blake3		wasm-rsa	
			Client 1	Client 2	Baseline	Client 1
Binary Before DCE	Functions	81	81	81	785	785
	Edges	269	269	269	9501	9501
	Reduction%	0.0	22.22	19.75	0.00	22.93
WASM-METADCE	Functions	81	63	65	785	605
	Edges	269	221	228	9501	8840
	Reduction%	0.0	22.22	19.75	0.00	22.93
WASMBRIDGE	Functions	65	47	49	776	596
	Edges	151	103	110	3652	2991
	Reduction%	19.75	41.98	39.51	1.15	24.08

Table 5. Comparison of Dead Code Elimination performed by WASM-METADCE and WASMBRIDGE. Reduction% shows the reduction in the number of functions after dead-code elimination.

RQ4: How does our approach compare to other state-of-the-art call graph analysis tools?

We compare our approach with other state-of-the-art call graph analysis tools, WASM-OPT[41], WASSAIL[37] and WASMA[5]. Several tools were excluded from our evaluation: WASSILLY [31] timed out on each subject application while generating a call graph, STURDY [22] does not support analysis over binaries that interoperate with JavaScript (the framework throws an error), and WASMATI [6] does not support WebAssembly functions that return more than one result (multi-value proposal). Table 4 compares the different analysis tools and we discuss their call graph analyses below.

- *WASM-OPT* performs a naive call graph analysis. At every indirect call site, it assumes that all the functions in the function table may be called. This results in a call graph with up to 3.8× more edges than our approach (e.g., 1,903 edges vs. 504 for webcam-stream). While this is the most sound analysis possible, the resulting call graph is very imprecise.
- *WASMA* performs a closed-world type-based call graph analysis. At every indirect call site, it considers all the functions in the table with the same function type as the type annotation at the call site to be a potential target of this call site. As expected from the results of RQ1 and 2, we see a reduction in the number of edges compared to WASMA. Note that WASMA removes functions that do not have any incoming or outgoing edges, which leads to a reduction in the number of nodes for subjects like epqueue. Our analysis does not remove these functions. Additionally, WASMA fails to run on one subject.
- *WASSAIL* performs a closed-world type-based call graph analysis like WASMA but like us, does not remove functions that do not have any incoming or outgoing edges.

RQ5: What is the effect of the increased precision of our approach on a downstream analysis like dead-code elimination?

To demonstrate the potential impact of our analysis technique on downstream analyses and how the addition of a JavaScript client analysis performs, we perform dead-code elimination to specialize a Wasm binary to a JavaScript client. For two subject applications, blake3 and wasm-rsa, the dataset we obtain them from has mock JavaScript clients for these subjects. The JavaScript clients call a subset of exported functions from the JavaScript wrapper of the WebAssembly binary. Through manual inspection, we determine the set of WebAssembly exported functions that are called by the mock client. We pass this information to WASM-METADCE, an industry state-of-the-art WebAssembly dead-code elimination tool, which, given a WebAssembly binary and a list of exported functions that are called by a client, performs dead-code elimination and produces a smaller binary that has been specialized to that client. We also use the list of exported functions called by the client to perform dead-code elimination based on our call graph analysis. We then compare the number of

883 functions that have been removed by WASM-METADCE versus WASMBRIDGE in Table 5. As a baseline,
 884 we pass in all the exported functions to both tools, to measure dead code elimination for a client
 885 that uses all exported functions.

- 886 • *Function Reduction*: WASMBRIDGE identifies more functions to be dead than WASM-METADCE,
 887 across both subjects and their clients. For blake3, we achieve function reductions of 19.75%,
 888 41.98%, and 39.51% compared to metadce's 0%, 22.22%, and 19.75% respectively. The reduction
 889 difference is not as high in wasm-rsa, but we still outperform WASM-METADCE.
- 890 • *Edge Reduction*: WASMBRIDGE continues to produce more precise call graphs with fewer
 891 edges, which can enable further optimizations and analyses.
- 892 • *Baseline*: Even in the baseline case, where a mock client uses all exported functions of
 893 a binary, our analysis identifies dead code that WASM-METADCE misses entirely (19.75%
 894 reduction vs. 0% for blake3).

896 7 RELATED WORK

897 *Inferring Richer Types for Binary Analysis*. Lehmann et al. [27] present SNOWWHITE, a
 898 learning-based approach that recovers high-level function types for WebAssembly functions using
 899 DWARF information in binaries. Zhao et al. [43] present WASMHINT, which infers ‘Rust-like’
 900 function types for WebAssembly functions using semantic learning in combination with program
 901 slicing. We do not compare against these approaches since they infer much more complicated types
 902 like structs and arrays and don’t provide an estimation of the soundness of their type systems or
 903 how the inferred types affect program analysis. However, both systems do infer a ptr type.

904 There has also been interest in recovering high-level type information for x86 binaries. Lee et
 905 al. [25] recover high-level types for x86 binaries by inferring types based on how code operates
 906 on data in the binary, similar to our inference algorithm over WebAssembly, while Bosamiya et al.
 907 [4] use a deductive type reconstruction strategy. ElWazeer et al. [12] use a “best-effort” flow and
 908 context-insensitive pointer analysis to recover type and variable information from x86 binaries.

909 *Multi-Language Static Analysis*. We now summarize alternate approaches to multi-language
 910 static analysis for other language pairs. Lee et al. [26] analyze JNI interoperation between Java
 911 and C by generating semantic summaries for C functions called by Java code. These summaries
 912 are translated into Java and calls to C functions are replaced with their Java summaries. Liang et
 913 al. [29] use LLVM IR as a common intermediate from multiple languages and detect bugs using
 914 symbolic execution coupled with Z3. Fornaia et al. [14] translates C, C++ and Java code into LLVM.
 915 Roth[34] reuses existing single-language analyses and combines them to support multiple languages
 916 by developing architecture that can be used to combine analyses. Monat et al. [30] implement
 917 an extensive C-Python abstract interpretation framework using transfer functions to convert C
 918 values to Python and vice-versa. Buro et al. [7] provide a general, theoretical framework combining
 919 abstract interpretations of different languages. Furr and Foster [17] also perform multilingual type
 920 inference for C code using client OCaml and Java types for checking the type-safety of FFI calls.

921 *WebAssembly Static Program Analysis*. There are several static analysis frameworks for We-
 922 bAssembly. We compare against WASSAIL[36], WASMA[5], WASM-OPT[41] and WASM-METADCE[41].
 923 The latter two are part of the Binaryen toolchain. WASSILY [31] and STURDY [22] improve the
 924 precision of indirect calls by relying on a better value analysis to determine the index into the
 925 function table. They are both abstract interpretation frameworks. We do not compare against
 926 them because the former times out on all our subjects and the latter does not support analysis of
 927 binaries that interoperate with JavaScript. WASMATI [6] uses code property graphs to statically
 928 detect vulnerabilities in WebAssembly binaries and use a type-based call graph analysis.

Refinement Type Systems. Refinement types have long been used to enhance type systems with logical predicates that constrain the set of values described by a type [15, 16, 32, 39, 42], allowing programmers to express more precise program properties and enabling program verification. The most closely related work is WasmPrecheck [18], a richer type system for WebAssembly that uses *indexed types* to express static constraints that enable safe removal of dynamic checks for type and memory safety. WasmPrecheck supports general constraints on index terms, while we support only singleton refinements and refining i32 to $\text{ptr}(a)$ or $\text{num}(n)$. Another point of comparison is that both WasmPrecheck and our work shows how to embed a Wasm module into a refined type system. However, Geller et al. give only a naive embedding — they show that a Wasm module can be (automatically) embedded into WashPrecheck by replacing all type annotations in the Wasm module with indexed types that have no constraints. This does not enhance the performance of the program, while our analysis considers information from interoperation with JavaScript to discover refinement types that in turn lead to improved static analysis.

8 CONCLUSION AND FUTURE WORK

We have presented the first multi-language static analysis for WebAssembly and JavaScript. Our work shows how a small refinement to WebAssembly’s type system, supported by value and pointer analysis, improves the precision of a call graph analysis by reducing the number of estimated call targets by an average of 17% at an average of 35.63% of indirect call sites. We also estimate the effect of this improved precision on downstream dead-code elimination analyses and outperform an industry state-of-the-art tool. However, our analysis only supports WebAssembly version 1.0 and JavaScript wrapper files generated by the Rust compiler. We plan to extend our analysis to support other versions of WebAssembly as well as binaries and wrappers produced by other toolchains.

Supporting WebAssembly 2.0 and 3.0. WebAssembly’s first version update introduced vector instructions, bulk memory instructions, multi-value results, non-trapping conversions and sign-extension instructions. We already support multi-value results and do not anticipate trouble when including these new instructions. WebAssembly 3.0 introduced Garbage Collection (GC), which has aggregate types like structures and arrays and linear memory is now 32- and 64-bit addressable. To support the GC proposal, we would first refine i64 types to have a similar refinement to i32 and extend the type system to include aggregate types. Since our refinement introduces subtype constraints between pointers and integers, we would still see precision improvements.

JavaScript Wrapper Analysis. Currently, our analysis only supports a CodeQL analysis of JavaScript wrappers generated by `wasm-bindgen`. However, this is not the only option to support an FFI between WebAssembly and JavaScript. C/C++ library developers typically use Emscripten to compile code to WebAssembly and JavaScript. While Emscripten provides several interoperation mechanisms between the two languages [13], Embind is the most popular mechanism and employs the same lowering and lifting as the wrappers produced by `wasm-bindgen` do.

Improvements to Precision. Our analysis discovers refinement types that are flow-insensitive, context-insensitive and intra-procedural. Additionally, the domain of values is fairly simple — a flat lattice of natural numbers. Considering the relative simplicity of our analysis we plan on employing several traditional techniques to increase precision of the analysis. We also plan to extend our CodeQL analysis to the JavaScript client to determine the values being passed to WebAssembly or being stored into the WebAssembly linear memory and the subset of exported functions being used by the client. The former would improve precision and the latter would enable a dead-code elimination analysis that specializes a WebAssembly binary to a JavaScript client.

981 REFERENCES

- 982 [1] 2023. *CodeQL: A Semantic Code Analysis Engine*. <https://codeql.github.com/>
- 983 [2] Ioannis Agadakos, Di Jin, David Williams-King, Vasileios P. Kemerlis, and Georgios Portokalidis. 2019. Nibbler: 984 debloating binary shared libraries. In *Proceedings of the 35th Annual Computer Security Applications Conference* 985 (San Juan, Puerto Rico, USA) (ACSAC '19). Association for Computing Machinery, New York, NY, USA, 70–83. 986 <https://doi.org/10.1145/3359789.3359823>
- 987 [3] Saed Alrabae, Mourad Debbabi, Paria Shirani, Lingyu Wang, Amr Youssef, Ashkan Rahimian, Lina Nouh, Djedjiga 988 Mouheb, He Huang, and Aiman Hanna. 2020. *Binary code fingerprinting for cybersecurity: Application to malicious code* 989 *fingerprinting*. Springer.
- 990 [4] Jay Bosamiya, Maverick Woo, and Bryan Parno. 2025. TReX: practical type reconstruction for binary code. In *Proceedings* 991 *of the 34th USENIX Conference on Security Symposium* (Seattle, WA, USA) (SEC '25). USENIX Association, USA, Article 992 354, 19 pages.
- 993 [5] Florian Breitfelder, Tobias Roth, Lars Baumgärtner, and Mira Mezini. 2023. WasmA: A Static WebAssembly Analysis 994 Framework for Everyone. In *2023 IEEE International Conference on Software Analysis, Evolution and Reengineering* 995 (SANER). 753–757. <https://doi.org/10.1109/SANER56733.2023.00085>
- 996 [6] Tiago Brito, Pedro Lopes, Nuno Santos, and José Fragoso Santos. [n. d.]. Wasmati: An Efficient Static Vulnerability 997 Scanner for WebAssembly. ([n. d.]), 102745. <https://doi.org/10.1016/j.cose.2022.102745>
- 998 [7] Samuele Buro, Roy L Crole, and Isabella Mastroeni. 2020. On multi-language abstraction: Towards a static analysis of 999 multi-language programs. In *Static Analysis: 27th International Symposium, SAS 2020, Virtual Event, November 18–20, 2020, Proceedings 27*. Springer, 310–332.
- 999 [8] Chris Fallin. 2025. *Waffle: WebAssembly Analysis Framework for Lightweight Experimentation*. Bytecode Alliance. 999 <https://github.com/bytocodealliance/waffle>
- 1000 [9] Ravi Chugh, David Herman, and Ranjit Jhala. 2012. Dependent types for JavaScript. In *Proceedings of the ACM 1001 international conference on Object oriented programming systems languages and applications*. 587–606.
- 1002 [10] Ron Cytron, Jeanne Ferrante, Barry K Rosen, Mark N Wegman, and F Kenneth Zadeck. 1989. An efficient method of 1003 computing static single assignment form. In *Proceedings of the 16th ACM SIGPLAN-SIGACT symposium on Principles of 1004 programming languages*. 25–35.
- 1005 [11] Nicolaas Govert De Bruijn. 1972. Lambda calculus notation with nameless dummies, a tool for automatic formula 1006 manipulation, with application to the Church-Rosser theorem. In *Indagationes mathematicae (proceedings)*, Vol. 75. Elsevier, 381–392.
- 1007 [12] Khaled ElWazeer, Kapil Anand, Aparna Kotha, Matthew Smithson, and Rajeev Barua. 2013. Scalable variable and data 1008 type detection in a binary rewriter. In *Proceedings of the 34th ACM SIGPLAN conference on Programming language design and implementation*. 51–60.
- 1009 [13] Emcc Docs Interacting with Code [n. d.]. *Emscripten Docs: Connecting C++ and JavaScript*. https://emscripten.org/docs/porting/connecting_cpp_and_javascript/Interacting-with-code.html
- 1010 [14] Andrea Fornaia, Stefano Scafati, and Emiliano Tramontana. 2019. JSCAN: Designing an Easy to use LLVM-Based 1011 Static Analysis Framework. In *2019 IEEE 28th International Conference on Enabling Technologies: Infrastructure for 1012 Collaborative Enterprises (WETICE)*. 237–242. <https://doi.org/10.1109/WETICE.2019.00058>
- 1013 [15] Jeremy Freeman, Catalin Hritcu, Marco Gaboardi, Vincent Laporte, Sergey Firsov, Gregory Malecha, Dustin Swasey, 1014 Conor McBride Watt, and Benjamin C. Pierce. 2020. RefinedC: Automating the Foundational Verification of C 1015 Code with Refined Types. In *Proceedings of the ACM on Programming Languages (OOPSLA)*, Vol. 4. 1–29. <https://doi.org/10.1145/3428198>
- 1016 [16] Tim Freeman and Frank Pfenning. 1991. Refinement types for ML. In *Proceedings of the ACM SIGPLAN 1991 conference 1017 on Programming language design and implementation*. 268–277.
- 1018 [17] Michael Furr and Jeffrey S Foster. 2005. Checking type safety of foreign function calls. *ACM SIGPLAN Notices* 40, 6 1019 (2005), 62–72.
- 1020 [18] Adam T. Geller, Justin P. Frank, and William J. Bowman. 2024. Indexed Types for a Statically Safe WebAssembly. 1021 *Proceedings of the ACM on Programming Languages* 8, POPL (2024), 2395–2424. <https://doi.org/10.1145/3632922>
- 1022 [19] Jessica Gronski, Kenneth Knowles, Aaron Tomb, Stephen N Freund, and Cormac Flanagan. 2006. Sage: Hybrid checking 1023 for flexible specifications. In *Scheme and Functional Programming Workshop*, Vol. 6. 93–104.
- 1024 [20] Andreas Haas, Andreas Rossberg, Derek L Schuff, Ben L Titzer, Michael Holman, Dan Gohman, Luke Wagner, Alon 1025 Zakai, and JF Bastien. 2017. Bringing the Web up to Speed with WebAssembly. In *Proceedings of the 38th ACM SIGPLAN 1026 Conference on Programming Language Design and Implementation* (Barcelona, Spain) (PLDI 2017). ACM, New York, NY, USA, 185–200. <https://doi.org/10.1145/3062341.3062363>
- 1027 [21] Aaron Hilbig, Daniel Lehmann, and Michael Pradel. 2021. An Empirical Study of Real-World WebAssembly Binaries: 1028 Security, Languages, Use Cases. In *Proceedings of the Web Conference 2021* (Ljubljana, Slovenia) (WWW '21). Association 1029 for Computing Machinery, New York, NY, USA, 2696–2708. <https://doi.org/10.1145/3442381.3450138>

- 1030 [22] Sven Keidel and Sebastian Erdweg. 2019. Sound and reusable components for abstract interpretation. *Proc. ACM
1031 Program. Lang.* 3, OOPSLA, Article 176 (Oct. 2019), 28 pages. <https://doi.org/10.1145/3360602>
- 1032 [23] Johannes Kinder. 2010. *Static analysis of x86 executables*. Ph.D. Dissertation. Technische Universität Darmstadt.
- 1033 [24] Christopher Kruegel, William Robertson, Fredrik Valeur, and Giovanni Vigna. 2004. Static disassembly of obfuscated
1034 binaries. In *USENIX security Symposium*, Vol. 13. 18–18.
- 1035 [25] JongHyup Lee, Thanassis Avgerinos, and David Brumley. 2011. TIE: Principled reverse engineering of types in binary
1036 programs. (2011).
- 1037 [26] SungHo Lee, Hyogun Lee, and Sukyoung Ryu. 2020. Broadening horizons of multilingual static analysis: Semantic
1038 summary extraction from C code for JNI program analysis. In *Proceedings of the 35th IEEE/ACM International Conference
1039 on Automated Software Engineering*. 127–137.
- 1040 [27] Daniel Lehmann and Michael Pradel. 2022. Finding the Dwarf: Recovering Precise Types from WebAssembly Binaries.
1041 In *Proceedings of the 43rd ACM SIGPLAN International Conference on Programming Language Design and Implementation*
1042 (New York, NY, USA) (PLDI ’22). Association for Computing Machinery. <https://doi.org/10.1145/3519939.3523449>
- 1043 [28] Daniel Lehmann, Michelle Thalakkottur, Frank Tip, and Michael Pradel. 2023. That’s a Tough Call: Studying the
1044 Challenges of Call Graph Construction for WebAssembly. In *Proceedings of the 32nd ACM SIGSOFT International
1045 Symposium on Software Testing and Analysis (ISSTA ’23)*. Association for Computing Machinery, New York, NY, USA.
1046 <https://doi.org/10.1145/3597926.3598104>
- 1047 [29] Hongliang Liang, Lei Wang, Dongyang Wu, and Jiuyun Xu. 2016. MLSA: A static bugs analysis tool based on LLVM
1048 IR. In *2016 17th IEEE/ACIS International Conference on Software Engineering, Artificial Intelligence, Networking and
1049 Parallel/Distributed Computing (SNPD)*. 407–412. <https://doi.org/10.1109/SNPD.2016.7515932>
- 1050 [30] Raphaël Monat, Abdelraouf Ouadjaout, and Antoine Miné. 2021. A multilanguage static analysis of python programs
1051 with native C extensions. In *International Static Analysis Symposium*. Springer, 323–345.
- 1052 [31] Mattia Paccamuccio, Franco Raimondi, and Michele Loreti. 2024. Building Call Graph of WebAssembly Programs via
1053 Abstract Semantics. arXiv:2407.14527 [cs.SE] <https://arxiv.org/abs/2407.14527>
- 1054 [32] Patrick M. Rondon, Kohei Kawaguchi, and Ranjit Jhala. 2008. Liquid Types. In *Proceedings of the ACM SIGPLAN
1055 Conference on Programming Language Design and Implementation (PLDI)*. 159–169.
- 1056 [33] Patrick Maxim Rondon, Ming Kawaguchi, and Ranjit Jhala. 2010. Low-level liquid types. *ACM Sigplan Notices* 45, 1
1057 (2010), 131–144.
- 1058 [34] Tobias Roth. 2023. Reusing Single-Language Analyses for Static Analysis of Multi-language Programs. In *Companion
1059 Proceedings of the 2023 ACM SIGPLAN International Conference on Systems, Programming, Languages, and Applications:
1060 Software for Humanity*. 16–18.
- 1061 [35] Dawn Song, David Brumley, Heng Yin, Juan Caballero, Ivan Jager, Min Gyung Kang, Zhenkai Liang, James Newsome,
1062 Pongsin Poosankam, and Prateek Saxena. 2008. BitBlaze: A new approach to computer security via binary analysis.
1063 In *Information Systems Security: 4th International Conference, ICSS 2008, Hyderabad, India, December 16–20, 2008.
1064 Proceedings 4*. Springer, 1–25.
- 1065 [36] Quentin Stiévenart and Coen De Roover. 2021. Wassail: a WebAssembly Static Analysis Library (*ProWeb21*). <https:////2021.programming-conference.org/home/proweb-2021> Fifth International Workshop on Programming Technology
1066 for the Future Web.
- 1067 [37] Quentin Stiévenart and Coen De Roover. 2020. Compositional Information Flow Analysis for WebAssembly Programs.
1068 In *2020 IEEE 20th International Working Conference on Source Code Analysis and Manipulation (SCAM)*. 13–24. <https://doi.org/10.1109/SCAM51674.2020.00007>
- 1069 [38] Michelle Thalakkottur, Max Bernstein, Daniel Lehmann, Michael Pradel, and Frank Tip. 2026. An Empirical Study of
1070 WebAssembly Usage in Node.js. In *48rd IEEE/ACM International Conference on Software Engineering, ICSE 2026*.
- 1071 [39] Niki Vazou, Eric Seidel, and Ranjit Jhala. 2014. Refinement Types for Haskell. In *Proceedings of the ACM SIGPLAN
1072 International Conference on Functional Programming (ICFP)*. 269–282.
- 1073 [40] Andreas Rossberg (Ed.). 2022-03-31. *WebAssembly 1.0 Core Specification*. World Wide Web Consortium (W3C).
1074 <https://www.w3.org/TR/wasm-core-1/>
- 1075 [41] WebAssembly Community Group. 2025. Binaryen: Optimizer and compiler/toolchain library for WebAssembly.
1076 <https://github.com/WebAssembly/binaryen>.
- 1077 [42] Hongwei Xi and Frank Pfenning. 1999. Dependent Types in Practical Programming. In *Proceedings of the ACM SIGPLAN
1078 Symposium on Principles of Programming Languages (POPL)*. 214–227.
- 1079 [43] Kunsong Zhao, Zihao Li, Weimin Chen, Xiapu Luo, Ting Chen, Guozhu Meng, and Yajin Zhou. 2025. Recasting
1080 Type Hints from WebAssembly Contracts. *Proc. ACM Softw. Eng.* 2, FSE, Article FSE118 (June 2025), 24 pages.
1081 <https://doi.org/10.1145/3729388>

1079 Appendices

1081 A INSTRUCTION TYPING FOR THE REFINEMENT TYPE SYSTEM

1083 $C^r ::= \{\text{func } tf_r^*, \text{ local } \tau_r^*, \text{ global } \tau_r^*, \text{ table } n?, \text{ memory } n?, \text{ label } (\tau_r^*, \Psi)^*, \text{ return } (\tau_r^*, \Psi)^?\}$

1085 Typing WebAssembly Instructions with Refinement Types

$R; C^r \vdash e^* : tf_r$

$$\frac{R = \{\text{inst } C^r^*, \text{tab } n^*, \text{mem } \Psi^*\} \quad (R \vdash \text{inst} : C^r)^* \quad (n \leq |cl|^*)^* \quad \vdash s.\text{mem} : R_\Psi}{\vdash \{\text{inst } inst^*, \text{tab } (cl^*)^*, \text{mem } (b^*)^*\} : R}$$

$$\frac{\forall a \in 0 \dots \text{size}(s.\text{mem}) \quad s.\text{mem}(a) = R_\Psi(a)}{\vdash s.\text{mem} : R_\Psi}$$

$$\begin{array}{c} \varsigma \text{ fresh} \quad \varsigma; \cdot; C^\alpha \vdash ex^* \text{ func } tf_w \text{ local } \tau_w^\ell e^* : \alpha^m \rightarrow \alpha^n; S'; \varsigma' \\ \rho = \text{solve}(S') \end{array}$$

$$\begin{array}{cccc} \tau_r^m = \rho[\alpha^m] & \tau_r^n = \rho[\alpha^n] & \Sigma_{\text{pre}} = \rho[C] & \Sigma_{\text{post}} = \rho[C'] \\ & & \exists \Psi_{\text{pre}}, \Psi_{\text{post}}, \exists \eta_{\text{pre}}, \eta_{\text{post}}. \end{array}$$

$$\frac{\text{Bij}(\eta_{\text{pre}}) \wedge \text{Bij}(\eta_{\text{post}}) \wedge \eta_{\text{pre}}(\Sigma_{\text{pre}}) = \Psi_{\text{pre}} \wedge \eta_{\text{post}}(\Sigma_{\text{post}}) = \Psi_{\text{post}} \wedge \eta_{\text{post}} \supseteq \eta_{\text{pre}}}{tf_r = \tau_r^m, \Psi_{\text{pre}} \rightarrow \tau_r^n, \Psi_{\text{post}}}$$

$$C^\alpha \vdash ex^* \text{ func } tf_w \text{ local } \tau_w^\ell e^* : ex^* tf_r; \rho$$

$$\alpha_w^g \text{ fresh}$$

$$C^\alpha = \{\text{func } tf_w^*, \text{global } \alpha_w^g, \text{table } n?, \text{memory } n?\}$$

$$(C^\alpha \vdash ex^* \text{ func } tf_w \text{ local } \tau_w^l e^* : ex^* tf_r; \rho)^*$$

$$C^r = \{\text{func } tf_r^*, \text{global } \tau_r^g, \text{table } n?, \text{memory } n?\}$$

$$\frac{(\text{module } f^* \text{ global } \tau_w^g \text{ tab? mem?}) \rightsquigarrow (\text{module } f^* \text{ global } \tau_r^g \text{ tab? mem?})}{(R; C^r \vdash f : ex^* tf_r)^*}$$

$$R \vdash \text{module } f^* \text{ global } \tau_w^g \text{ tab? mem?} \quad \text{MODULE}$$

$$tf_r = \tau_r^m, \Psi_{\text{pre}} \rightarrow \tau_r^n, \Psi_{\text{post}}$$

$$R, C^r, \text{local}(\tau_r^m, \tau_r^\ell), \text{label}(\tau_r^n, \Psi_{\text{post}}), \text{return}(\tau_r^n, \Psi_{\text{post}}) \vdash e^* : \epsilon \rightarrow \tau_r^n \quad \text{FUNCTION}$$

$$R; C^r \vdash ex^* \text{ func } tf_r \text{ local } \tau_r^\ell e^* : ex^* tf_r$$

$$tf_r = \tau_r^m, \Psi_{\text{pre}} \rightarrow \tau_r^n, \Psi_{\text{post}} \quad R_\Psi = \Psi_{\text{pre}}$$

$$R; C^r, \text{label}(\tau_r^n, \Psi_{\text{post}}) \vdash e^* : \tau_r^m, \Psi_{\text{pre}} \rightarrow (\tau_r')^n, \Psi'_{\text{post}}$$

$$(\tau_r' <: \tau_r)^n \quad \Psi'_{\text{post}} <: \Psi_{\text{post}}$$

$$R; C^r \vdash \text{block } tf_r e^* \text{ end} : tf_r \quad \text{BLOCK}$$

$$tf_r = \tau_r^m, \Psi_{\text{pre}} \rightarrow \tau_r^n, \Psi_{\text{post}} \quad R_\Psi = \Psi_{\text{pre}}$$

$$R; C^r, \text{label}(\tau_r^m, \Psi_{\text{pre}}) \vdash e^* : \tau_r^m, \Psi_{\text{pre}} \rightarrow (\tau_r')^n, \Psi'_{\text{post}}$$

$$(\tau_r' <: \tau_r)^n \quad \Psi'_{\text{post}} <: \Psi_{\text{post}}$$

$$R; C^r \vdash \text{loop } tf_r e^* \text{ end} : tf_r \quad \text{LOOP}$$

Fig. 7. Typing Rules for the Refinement Type System

```

1128 def load_and_extend( $a, \tau, tp\_sz^?, \Psi$ ) :=  

1129   let  $\tau_r(n_1) = \Psi(a)$   

1130   let  $\tau_r(n_2) = \Psi(a - 4)$   

1131   if  $tp$  given then let  $N = |tp|$  else let  $N = |\tau|$   

1132   if  $wasm\_type(\tau_r(n_1)) == \tau$  then let  $\tau_r(n) = \tau_r(n_1)$   

1133   if  $|\tau_r(n_1)| < N$  then let  $\tau_r(n) = \tau(\text{val}(\text{bits}(n_2) \text{ bits}(n_1)))$   

1134   if  $|\tau_r(n_1)| > N$  then let  $\tau_r(n) = \tau(\text{val}(\text{bits}(n_1)_{0..N}))$   

1135   if  $(\tau_r(n_1) = (\text{i32}(n_1) \vee \text{ptr}(n_1) \vee \text{num}(n_1)) \wedge \tau = \text{f32})$  then let  $\tau_r(n) = \text{f32}(\text{bits}(n_1 \text{ as } \text{f32}))$   

1136   if  $(\tau_r(n_1) = \text{f32}(n_1) \wedge \tau = (\text{i32} \vee \text{ptr} \vee \text{num}))$  then let  $\tau_r(n) = \text{i32}(\text{f32}(n_1) \text{ as } \text{i32})$   

1137  

1138 if  $tp$  given  

1139   then return  $\tau_r(\text{extend\_sx}_{N, |\tau_w|}(n))$   

1140   else return  $\tau_r(n)$   

1141  

1142  

1143  

1144  $C_{\text{memory}}^r = n \quad \tau_r^1 <: \text{ptr}(c) \quad \tau_r^2 = \text{load\_and\_extend}(c + o, \tau, (tp\_sx)^?, R_\Psi)$   

1145  $2^a \leq (|tp| <)^\? |\tau_r^2| \quad (tp\_sz)^\? = \epsilon \vee \tau_r^2 = \text{im}$   

1146  $\frac{}{R; C^r \vdash \tau.\text{load}(tp\_sx)^\? a o : \tau_r^1 \rightarrow \tau_r^2}$  LOAD  

1147  

1148  

1149 def wrap_and_store( $p, \tau_r(n), \tau, tp^?, R$ ) :=  

1150   if  $tp$  then let  $N = |tp|$  else let  $N = |\tau|$   

1151   if  $|\tau| == 64$  then  $R_\Psi[p \mapsto \text{i32}(\text{val}(\text{bits}(n)_{0..32})) \wedge (p - 4) \mapsto \text{i32}(\text{val}(\text{bits}(n)_{32..64}))]$   

1152   if  $|\tau_r(n)| > N$  then  

1153      $R_\Psi[p \mapsto \tau_r(\text{wrap}_{|\tau|, N}(n))]$   

1154   else  $R_\Psi[p \mapsto \tau_r(n)]$   

1155  

1156  

1157  $C_{\text{memory}}^r = n \quad 2^a \leq (|tp| <)^\? |\tau_r^2| \quad tp^? = \epsilon \vee \tau_r^2 = \text{im}$   

1158  $\tau_r^1 <: \text{ptr32}(c) \quad \text{wrap\_and\_store}(c + o, \tau_r^2, \tau_w, tp^?, R) \quad \tau_r^2 <: \tau(\top)$   

1159  $\frac{}{R, C^r \vdash \tau.\text{store } tp^? a o : \tau_r^1 \tau_r^2 \rightarrow \epsilon}$  STORE  

1160  

1161  

1162  $C_{\text{memory}}^r = n \quad \text{num32}(n_1) <: \tau_r^1 \quad \text{num32}(n_2) <: \tau_r^2$   

1163  $\frac{}{R; C^r \vdash \text{grow\_memory} : \tau_r^1 \rightarrow \tau_r^2}$  GROW-MEMORY  

1164  

1165  

1166  $C_{\text{memory}}^r = n \quad \text{num32}(n_1) <: \tau_r \quad |R_\Psi| = n$   

1167  $\frac{}{R; C^r \vdash \text{current\_memory} : \epsilon \rightarrow \tau_r}$  CURRENT-MEMORY  

1168  

1169  

1170  $C_{\text{func}}^r(i) = (\tau_r')^m, \Psi'_{\text{pre}} \rightarrow (\tau_r')^n, \Psi'_{\text{post}} \quad tf_r = \tau_r^m, \Psi_{\text{pre}} \rightarrow \tau_r^n, \Psi_{\text{post}}$   

1171  $(\tau_r <: \tau_r')^m \quad (\tau_r' <: \tau_r)^n \quad \Psi'_{\text{pre}} <: \Psi_{\text{pre}} \quad \Psi_{\text{post}} <: \Psi'_{\text{post}}$   

1172  $\frac{}{R; C^r \vdash \text{call } i \ tf_r : tf_r}$  CALL  

1173  

1174  

1175  $tf_r = \tau_r^m, \Psi_{\text{pre}} \rightarrow \tau_r^n, \Psi_{\text{post}} \quad C_{\text{table}}^r = n \quad \text{num32}(c) <: \tau_r^1$   

1176  $R; C^r \vdash \text{call\_indirect } tf_r : \tau_r^m, \tau_r^1 \rightarrow \tau_r^n, \Psi_{\text{post}}$  CALL-INDIRECT  

1177  

1178 Fig. 8. Typing Rules for Refinement Type System continued.  

1179

```

$$\begin{array}{c}
1177 \quad \frac{}{R; C^r \vdash \tau.\mathbf{const} c : \epsilon \rightarrow \tau(c)} \text{CONST} \qquad \frac{\tau_r^2 = \mathbf{unop}(\tau_r^1)}{R; C^r \vdash \tau.\mathbf{unop} : \tau_r^1 \rightarrow \tau_r^2} \text{UNOP} \\
1178 \\
1179 \\
1180 \quad \frac{\tau_r^3 = \mathbf{binop}(\tau_r^1, \tau_r^2) \quad \tau_r^3 <: \tau(\top)}{R; C^r \vdash \tau.\mathbf{binop} : \tau_r^1 \tau_r^2 \rightarrow \tau_r^3} \text{BINOP} \qquad \frac{\tau_r^2 = \mathbf{testop}(\tau_r^1) \quad \tau_r^2 <: \mathbf{num}(\top)}{R; C^r \vdash \tau.\mathbf{testop} : \tau_r^1 \rightarrow \tau_r^2} \text{TESTOP} \\
1181 \\
1182 \\
1183 \\
1184 \quad \frac{\tau_r^3 = \mathbf{relop}(\tau_r^1, \tau_r^2) \quad \tau_r^3 <: \mathbf{num}(\top)}{R; C^r \vdash \tau.\mathbf{relop} : \tau_r^1 \tau_r^2 \rightarrow \tau_r^3} \text{RELOP} \qquad \frac{}{R; C^r \vdash \mathbf{unreachable} : \tau_r^n \rightarrow \tau_r^m} \text{UNREACHABLE} \\
1185 \\
1186 \\
1187 \quad \frac{}{R; C^r \vdash \mathbf{nop} : \epsilon \rightarrow \epsilon} \text{NOP} \qquad \frac{}{R; C^r \vdash \mathbf{drop} : \tau_r \rightarrow \epsilon} \text{DROP} \\
1188 \\
1189 \\
1190 \quad \frac{}{R; C^r \vdash \mathbf{select} : \tau_r^1 \tau_r^2 \mathbf{num32}(n) \rightarrow \tau_r^3} \text{SELECT} \\
1191 \\
1192 \quad \frac{\begin{array}{c} tf_r = \tau_r^m, \Psi_{\text{pre}} \rightarrow \tau_r^n, \Psi_{\text{post}} \\ R_\Psi = \Psi_{\text{pre}} \\ R; C^r, \text{label}(\tau_r^n) \vdash e_1^* : \tau_r^m, \Psi_{\text{pre}} \rightarrow (\tau_r')^n, \Psi'_{\text{post}} \\ R; C^r, \text{label}(\tau_r^n) \vdash e_2^* : \tau_r^m, \Psi_{\text{pre}} \rightarrow (\hat{\tau}_r'')^n, \Psi''_{\text{post}} \\ (\tau_r' <: \tau_r)^n \quad (\hat{\tau}_r'' <: \tau_r)^n \quad \Psi'_{\text{post}} <: \Psi_{\text{post}} \quad \Psi''_{\text{post}} <: \Psi_{\text{post}} \end{array}}{R; C^r \vdash \mathbf{if} \ tf_r \ \mathbf{then} \ e_1^* \ \mathbf{else} \ e_2^* \ \mathbf{end} : tf_r} \text{IF-ELSE} \\
1193 \\
1194 \\
1195 \\
1196 \\
1197 \\
1198 \quad \frac{C_{\text{label}}^r(i) = (\tau_r')^m, \Psi' \quad (\tau_r <: \tau_r')^m \quad R_\Psi <: \Psi'}{R; C^r \vdash \mathbf{br} \ i : \tau_r^* \tau_r^m \rightarrow \tau_r^*} \text{BR} \\
1199 \\
1200 \\
1201 \\
1202 \quad \frac{\begin{array}{c} C_{\text{label}}^r(i) = (\tau_r')^m, \Psi' \quad \mathbf{num32}(n) <: \tau_r^1 \\ (\tau_r <: \tau_r')^m \quad R_\Psi <: \Psi' \end{array}}{R; C^r \vdash \mathbf{br_if} \ i : \tau_r^m \tau_r^1 \rightarrow \tau_r^m} \text{BR-IF} \qquad \frac{\begin{array}{c} C_{\text{label}}^r(i) = ((\tau_r')^m, \Psi)^+ \\ ((\tau_r <: \tau_r')^m)^+ \quad (R_\Psi <: \Psi)^+ \end{array}}{R; C^r \vdash \mathbf{br_table} \ i^+ : \tau_r^* \tau_r^m \tau_r^1 \rightarrow \tau_r^*} \text{BR-TABLE} \\
1203 \\
1204 \\
1205 \\
1206 \quad \frac{\begin{array}{c} C_{\text{label}}^r(i) = (\tau_r')^m, \Psi \\ (\tau_r <: \tau_r')^m \quad R_\Psi <: \Psi \end{array}}{R; C^r \vdash \mathbf{return} : \tau_r^* \tau_r^m \rightarrow \tau_r^*} \text{RETURN} \qquad \frac{C_{\text{local}}^r(i) = \tau_r}{R; C^r \vdash \mathbf{get_local} \ i : \epsilon \rightarrow \tau_r} \text{GET-LOCAL} \\
1207 \\
1208 \\
1209 \\
1210 \\
1211 \quad \frac{C_{\text{local}}^r(i) = \tau_r^2 \quad \tau_r^1 <: \tau_r^2}{R; C^r \vdash \mathbf{set_local} \ i : \tau_r^1 \rightarrow \epsilon} \text{SET-LOCAL} \qquad \frac{C_{\text{local}}^r(i) = \tau_r^2 \quad \tau_r^1 <: \tau_r^2}{R; C^r \vdash \mathbf{tee_local} \ i : \tau_r^1 \rightarrow \tau_r^1} \text{TEE-LOCAL} \\
1212 \\
1213 \\
1214 \quad \frac{C_{\text{global}}^r(i) = \tau_r}{R; C^r \vdash \mathbf{get_global} \ i : \epsilon \rightarrow \tau_r} \text{GET-GLOBAL} \qquad \frac{C_{\text{global}}^r(i) = \tau_r^2 \quad \tau_r^1 <: \tau_r^2}{R; C^r \vdash \mathbf{set_global} \ i : \tau_r^1 \rightarrow \epsilon} \text{SET-GLOBAL} \\
1215 \\
1216 \\
1217 \quad \frac{}{R; C^r \vdash \epsilon : \epsilon \rightarrow \epsilon} \text{EMPTYSTACK} \qquad \frac{\begin{array}{c} R; C^r \vdash e_1^* : \tau_r^a \rightarrow \tau_r^b \quad R; C^r \vdash e_2 : \tau_r^b \rightarrow \tau_r^c \\ R; C^r \vdash e_1^* \ e_2 : \tau_r^a \rightarrow \tau_r^c \end{array}}{R; C^r \vdash e_1^* \ e_2 : \tau_r^a \rightarrow \tau_r^c} \text{SEQUENCING} \\
1218 \\
1219 \\
1220 \\
1221 \quad \frac{R; C^r \vdash e^* : \tau_r^b \rightarrow \tau_r^c}{R; C^r \vdash e^* : \tau_r^a \ \tau_r^b \rightarrow \tau_r^a \ \tau_r^c} \text{TOPOFSTACK} \qquad \frac{}{R; C^r \vdash \mathbf{trap} : tf_r} \text{TRAP} \\
1222 \\
1223 \\
1224
\end{array}$$

Fig. 9. Typing Rules for the Refinement Type System continued.

1226 **B TYPE SAFETY PROOF**

1227 THEOREM 3. **Progress:** If $R; C^r \vdash_i \hat{e}^* : \tau_r^1 \cdots \tau_r^m \rightarrow \tau_r^* \wedge (v : \tau_r)^m \wedge \vdash_i s : R$ then $e^* = v^*$ or
 1228 $e^* = \text{trap}$ or $s; v^*; e^* \hookrightarrow s'; v'^*; e'^*$

1229 PROOF. Proof by induction of the derivation of $R; C^r \vdash_i \hat{e}^* : \tau_r^1 \cdots \tau_r^m \rightarrow \tau_r^*$
 1230 CASE LOAD:

$$\frac{\begin{array}{c} C_{\text{memory}}^r = n \quad \tau_r^1 \lessdot \text{ptr}(c) \quad \tau_r^2 = \text{load_and_extend}(c + o, \tau_w, (tp_sx)^?, R_\Psi) \\ 2^a \leq (|tp| <)^? |\tau_r^2| \quad (tp_sz)^? = \epsilon \vee \tau_r^2 = \text{im} \end{array}}{R; C^r \vdash \tau.\text{load}(tp_sx)^? a o : \tau_r^1 \rightarrow \tau_r^2} \text{LOAD}$$

1231 The load instruction has to take a step since only end instructions signify the end of a function. We
 1232 now case on the shape of a well typed stack. The only possible case is that the load instruction has
 1233 a $\text{ptr}(a)$ on the stack, since this is a premise of our typing rule. The operational semantics of the
 1234 load instruction matches this stack shape. If the side conditions are not met, it steps to a trap. \square

1235 THEOREM 4. **Preservation:** If $R; C^r \vdash_i \hat{e}^* : \tau_r^1 \cdots \tau_r^m \rightarrow \tau_r^* \wedge \vdash_i (v : \tau_r)^m \wedge \vdash_i s : R \wedge s, v^*, e^* \hookrightarrow s', v'^*, e'^*$ then $\exists R'$ such that $\text{dom}(R_\Psi) \subseteq \text{dom}(R'_\Psi) \wedge \vdash_i s' : R' \wedge R' ; C^r \vdash_i e'^* : \tau_r^*$

1236 PROOF. Proof by induction of the derivation of $R; C^r \vdash_i \hat{e}^* : \tau_r^1 \cdots \tau_r^m \rightarrow \tau_r^*$
 1237 CASE LOAD:

$$\frac{\begin{array}{c} C_{\text{memory}}^r = n \quad \tau_r^1 \lessdot \text{ptr}(c) \quad \tau_r^2 = \text{load_and_extend}(c + o, \tau_w, (tp_sx)^?, R_\Psi) \\ 2^a \leq (|tp| <)^? |\tau_r^2| \quad (tp_sz)^? = \epsilon \vee \tau_r^2 = \text{im} \end{array}}{R; C^r \vdash \tau.\text{load}(tp_sx)^? a o : \tau_r^1 \rightarrow \tau_r^2} \text{LOAD}$$

1238 From our typing rule for the load instruction we have that it expects τ_r^1 on the stack and pushes τ_r^2
 1239 on the stack. From the operational semantics of the load instruction we know that there are three
 1240 possible cases:

1241 **CASE 1:** $s; \text{ptr.const}(c); \tau.\text{load } a o \hookrightarrow \tau.\text{const}(b^*)$ if $s_{\text{mem}}(i, c + o, |\tau|) = b^*$

1242 From $\vdash_i s : R$, we know that $\vdash s_{\text{mem}} : R_\Psi$, from which we know that $s_{\text{mem}}(c + o) : R_\Psi(c + o)$.
 1243 Since our memory typing R_Ψ maps addresses to types of size 32, we have to now case on the data
 1244 expected after the load instruction. This is annotated on the load instruction itself as τ .

1245 CASE 1.1: $\tau.\text{load } a o \wedge R_\Psi[(c + o) \mapsto \tau_{r_{32}}(n)]$

1246 This is the case where the data in memory is of the type that the load expects to produce
 1247 on the stack. We now have to show that $b^* : \tau(n)$. From the typing rule, we have that load
 1248 expects τ_r^2 . On inspecting the load_and_extend function, we see that for the case where the data in
 1249 memory is the same type as is expected on the stack, $\tau_r^2 = \tau_r(n_1)$ where $\tau_r(n_1) = \Psi(c + o)$. Since
 1250 $s_{\text{mem}}(c + o) : R_\Psi(c + o)$, we know that $s_{\text{mem}}(c + o) : \tau_r^2$ and since $s_{\text{mem}}(c + o) = b^*$, $b^* : \tau_r^2$.

1251 CASE 1.2: $i64.\text{load } a o \wedge R_\Psi[(c + o) \mapsto \tau_{r_{32}}(n_1), (c + o - 4) \mapsto \tau_{r_{32}}(n_2)]$

1252 In this case, the load expects a i64 on the stack and the memory has $\tau_{r_{32}}(n_1)$ at $(c + o)$ and
 1253 $\tau_{r_{32}}(n_2)$ at $(c + o - 4)$. We now have to show that $b^* : i64(n)$. Note that b^* is the byte sequence
 1254 from $s_{\text{mem}}[c + o : 8]$, as specified in the detailed Wasm specification [40]. From the typing rule,
 1255 we know that load expects τ_r^2 on the stack. On inspecting the load_and_extend function, we see
 1256 that this case is the one that matches $|\tau_r(n_1)| < N$, where $N = |\tau| = 64$. We see then that
 1257 $\tau_r^2 = \tau(\text{val}(\text{bit}(n_1), \text{bit}(n_2)))$. n_1 and n_2 are the values in memory at $c + o$ and $c + o - 4$. Since
 1258 $s_{\text{mem}}(c + o) : R_\Psi(c + o) \wedge s_{\text{mem}}(c + o - 4) : R_\Psi(c + o - 4)$, we know that $s_{\text{mem}}(c + o : 8) : \tau_r^2$ and since
 1259 $s_{\text{mem}}(c + o) = b^*$, $b^* : \tau_r^2$.

1275 CASE 1.3: $f64.load\ a\ o \wedge R_\Psi[(c + o) \mapsto \tau_{r_{32}}(n_1), (c + o - 4) \mapsto \tau_{r_{32}}(n_2)]$

1276 In this case, the load expects a $f64$ on the stack and the memory has $\tau_{r_{32}}(n_1)$ at $(c + o)$ and $\tau_{r_{32}}(n_2)$ at $(c + o - 4)$. We now have to show that $b^* : f64(n)$. The proof proceeds exactly as the case above.

1278 CASE 1.4: $ptr.load\ a\ o \wedge R_\Psi[(c + o) \mapsto f32(n)]$

1279 CASE 1.5: $num.load\ a\ o \wedge R_\Psi[(c + o) \mapsto f32(n)]$

1280 CASE 1.6: $i32.load\ a\ o \wedge R_\Psi[(c + o) \mapsto f32(n)]$

1281 In this case, the load expects a $i32$ on the stack and Ψ has $f32(n)$ at $c + o$. We now have to show that $b^* : f32(n)$. From the typing rule, we have that load expects τ_r^2 . On inspecting the `load_and_extend` function, we see that for this case, $\tau_r^2 = f32(\text{bits}(n_1 \text{ as } f32))$ where $\tau_r(n_1) = \Psi(c + o)$. Since $s_{mem}(c + o) : R_\Psi(c + o)$, we know that $s_{mem}(c + o) : \tau_r^2$ and since $s_{mem}(c + o) = b^*$, $b^* : \tau_r^2$.

1286 CASE 1.7: $f32.load\ a\ o \wedge R_\Psi[(c + o) \mapsto ptr(n)]$

1287 CASE 1.8: $f32.load\ a\ o \wedge R_\Psi[(c + o) \mapsto num(n)]$

1288 CASE 1.9: $f32.load\ a\ o \wedge R_\Psi[(c + o) \mapsto i32(n)]$

1289 In this case, the load expects a $f32$ on the stack and Ψ has $i32(n)$ at $c + o$. We now have to show that $b^* : i32(n)$. From the typing rule, we have that load expects τ_r^2 . On inspecting the `load_and_extend` function, we see that for this case, $\tau_r^2 = i32(f32(n_1) \text{ as } i32)$ where $\tau_r(n_1) = \Psi(c + o)$. Since $s_{mem}(c + o) : R_\Psi(c + o)$, we know that $s_{mem}(c + o) : \tau_r^2$ and since $s_{mem}(c + o) = b^*$, $b^* : \tau_r^2$.

1293 CASE 2: $s; ptr.const(c); \tau.load\ tp_sx\ a\ o \hookrightarrow \tau.const(b^*)$ if $s_{mem}(i, c + o, |tp|) = b^*$

1294 From $\vdash s : R$, we know that $\vdash s_{mem} : R_\Psi$, from which we know that $s_{mem}(c + o) : R_\Psi(c + o)$. Since 1295 our memory typing R_Ψ maps addresses to types of size 32, we have to now case on the data expected 1296 after the load instruction. The proof proceeds as above except with a small change introduced by 1297 the `tp_sx` annotation on the load. In the WebAssembly operational semantics [40], when a load 1298 has this packed type annotation, the bits are read from memory upto $|tp|$ and the value read from 1299 memory is size extended with the `extend_sxN,|τ|` function. Both of these cases are handled in the 1300 `load_and_extend` function, in case $|\tau_r(n_1)| > N$ and size extending function `extend_sxN,|τ|`. 1301

1302 CASE 3: $s; ptr.const(c); \tau.load\ (tp_sx)?\ a\ o \hookrightarrow \text{trap}$ otherwise

1303 The trap instruction is well typed. □

1306 C MEET AND JOIN OPERATIONS OVER THE INTEGER SUB-LATTICE

\sqcup	$i32(\top)$	$i32(n)$	$ptr(\top, \top)$	$ptr(l_0, n_0)$	$ptr(l_1, n_1)$	$ptr(\perp, \perp)$	$num(\top)$	$num(n)$	$num(\perp)$	$i32(\perp)$
$i32(\top)$	$i32(\top)$	$i32(\top)$	$i32(\top)$	$i32(\top)$	$i32(\top)$	$i32(\top)$	$i32(\top)$	$i32(\top)$	$i32(\top)$	$i32(\top)$
$i32(n')$	$i32(\top)$	$i32(n \sqcup n')$	$i32(n')$	$i32(n' \sqcup n_0)$	$i32(\top)$	$i32(n')$	$i32(\top)$	$i32(n' \sqcup n)$	$i32(n')$	$i32(n')$
$ptr(\top, \top)$	$i32(\top)$	$i32(n)$	$ptr(\top, \top)$	$ptr(\top, \top)$	$ptr(\top, \top)$	$ptr(\top, \top)$	$i32(\top)$	$i32(\top)$	$ptr(\top, \top)$	$ptr(\top, \top)$
$ptr(l_0, n'_0)$	$i32(\top)$	$i32(n \sqcup n'_0)$	$ptr(\top, \top)$	$ptr(l_0, n'_0 \sqcup n_0)$	$ptr(\top, \top)$	$ptr(l_0, n'_0)$	$i32(\top)$	$i32(\top)$	$i32(n'_0)$	$ptr(l_0, n'_0)$
$ptr(l_1, n'_1)$	$i32(\top)$	$i32(\top)$	$ptr(\top, \top)$	$ptr(\top, \top)$	$ptr(l_1, n_1 \sqcup n'_1)$	$ptr(l_1, n'_1)$	$i32(\top)$	$i32(\top)$	$ptr(l_1, n'_1)$	$i32(\top)$
$ptr(l_2, n'_2)$	$i32(\top)$	$i32(\top)$	$ptr(\top, \top)$	$ptr(\top, \top)$	$ptr(l_2, n'_2)$	$i32(\top)$	$i32(\top)$	$i32(\top)$	$ptr(l_2, n'_2)$	$i32(\top)$
$ptr(\perp, \perp)$	$i32(\top)$	$i32(\top)$	$ptr(\top, \top)$	$ptr(l_0, n_0)$	$ptr(l_1, n_1)$	$ptr(\perp, \perp)$	$i32(\top)$	$i32(\top)$	$ptr(\perp, \perp)$	$i32(\top)$
$num(\top)$	$i32(\top)$	$i32(n)$	$i32(\top)$	$i32(\top)$	$i32(\top)$	$i32(\top)$	$num(\top)$	$num(\top)$	$num(\top)$	$num(\top)$
$num(n')$	$i32(\top)$	$i32(n \sqcup n')$	$i32(\top)$	$i32(n_0 \sqcup n')$	$i32(\top)$	$i32(\top)$	$num(n')$	$num(n' \sqcup n)$	$num(n')$	$num(n')$
$num(\perp)$	$i32(\top)$	$i32(n)$	$i32(\top)$	$i32(n_0)$	$i32(\top)$	$i32(\top)$	$num(\top)$	$num(n)$	$num(\perp)$	$num(\perp)$
$i32(\perp)$	$i32(\top)$	$i32(n)$	$ptr(\top, \top)$	$ptr(l_0, n_0)$	$ptr(l_1, n_1)$	$ptr(\perp, \perp)$	$num(\top)$	$num(n)$	$num(\perp)$	$i32(\perp)$

1321 Table 6. Join \sqcup operation for the $i32$ sub-lattice.

	\sqcap	i32(\top)	i32(n)	ptr(\top, \top)	ptr(l_0, n_0)	ptr(l_1, n_1)	ptr(\perp, \perp)	num(\top)	num(n)	num(\perp)	i32(\perp)
1324	i32(\top)	i32(\top)	i32(n)	ptr(\top, \top)	ptr(l_0, n_0)	ptr(l_1, n_1)	ptr(\perp, \perp)	num(\top)	num(n)	num(\perp)	i32(\perp)
1325	i32(n')	i32(n')	i32($n \sqcap n'$)	ptr(\top, \top)	ptr($l_0, n_0 \sqcap n'$)	ptr(\perp, \perp)	ptr(\perp, \perp)	num(n')	num($n \sqcap n'$)	num(\perp)	i32(\perp)
1326	ptr(\top, \top)	ptr(\top, \top)	ptr(\top, \top)	ptr(\top, \top)	ptr(l_0, n_0)	ptr(l_1, n_1)	ptr(\perp, \perp)	i32(\perp)	i32(\perp)	i32(\perp)	i32(\perp)
1327	ptr(l_0, n'_0)	ptr(l_0, n'_0)	ptr($l_0, n'_0 \sqcap n$)	ptr(l_0, n'_0)	ptr($l_0, n_0 \sqcap n'_0$)	ptr(\perp, \perp)	ptr(\perp, \perp)	i32(\perp)	i32(\perp)	i32(\perp)	i32(\perp)
1328	ptr(l_1, n'_1)	ptr(l_1, n'_1)	ptr(\perp, \perp)	ptr(l_1, n'_1)	ptr(\perp, \perp)	ptr($l_1, n_1 \sqcap n'_1$)	ptr(\perp, \perp)	i32(\perp)	i32(\perp)	i32(\perp)	i32(\perp)
1329	ptr(l_2, n'_2)	ptr(l_2, n'_2)	ptr(\perp, \perp)	ptr(l_2, n'_2)	ptr(\perp, \perp)	ptr($l_3, 0$) [*]	ptr(\perp, \perp)	i32(\perp)	i32(\perp)	i32(\perp)	i32(\perp)
1330	ptr(\perp, \perp)	ptr(\perp, \perp)	ptr(\perp, \perp)	ptr(\perp, \perp)	ptr(\perp, \perp)	ptr(\perp, \perp)	ptr(\perp, \perp)	i32(\perp)	i32(\perp)	i32(\perp)	i32(\perp)
1331	num(\top)	num(\top)	num(\top)	i32(\perp)	i32(\perp)	i32(\perp)	i32(\perp)	num(\top)	num(n)	num(\perp)	i32(\perp)
1332	num(n')	num(n')	num($n \sqcap n'$)	i32(\perp)	i32(\perp)	i32(\perp)	i32(\perp)	num(n')	num($n \sqcap n'$)	num(\perp)	i32(\perp)
1333	num(\perp)	num(\perp)	num(\perp)	i32(\perp)	i32(\perp)	i32(\perp)	i32(\perp)	num(\perp)	num(\perp)	num(\perp)	i32(\perp)
1334	i32(\perp)	i32(\perp)	i32(\perp)	i32(\perp)	i32(\perp)	i32(\perp)	i32(\perp)	i32(\perp)	i32(\perp)	i32(\perp)	i32(\perp)

* l_3 fresh $\wedge l_1 = l_3 - n_1 \wedge l_2 = l_3 - n'_2$

Table 7. Meet \sqcap operation for the i32 sub-lattice.

C.1 Join of two pointers

The canonical example for joins in static analysis is the if-then-else expression. Let us suppose that the then branch of such an expression returns $\text{ptr}(l_1, n_1)$, and the else branch returns $\text{ptr}(l_2, n_2)$. What pointer does the if-then-else expression return? Since it could be either of the two, we do not presume to know what the pointer could be and instead return $\text{ptr}(\top, \top)$. Instead, if the then branch returned $\text{ptr}(l_1, n_1)$ and the else branch returned $\text{ptr}(l_2, n_2)$, we say that the if-then-else expression returns $\text{ptr}(l_1, n_1 \sqcap n_2)$.

C.2 Meet of two pointers

The canonical example for meet operations in static analysis is values during several loop iterations. Say, at loop iteration n , the type of a stack slot is $\text{ptr}(l_1, n_1)$, a pointer with a symbolic address, and at loop iteration $n + 1$, its type is $\text{ptr}(l_c, n_0)$, a constant pointer. Since it is impossible for a stack slot to have a symbolic and constant pointer, we return $\text{ptr}(\perp, \perp)$. On the other hand, say that at loop iteration n , the type of a stack slot is $\text{ptr}(l_1, n_1)$ and at loop iteration $n + 1$, its type is $\text{ptr}(l_2, n_2)$. This means that the type at this stack slot is both $\text{ptr}(l_1, n_1)$ and $\text{ptr}(l_2, n_2)$. We equate these two pointers to a third pointer $\text{ptr}(l_3, 0)$, where l_3 is a fresh symbolic location and $\text{ptr}(l_1, n_1) = \text{ptr}(l_3, 0) \wedge \text{ptr}(l_2, n_2) = \text{ptr}(l_3, 0)$. Pointers can be written as polynomials (since the offset is usually added into the base address) and we get that $l_1 + n_1 = l_3$ or that, $l_1 = l_3 - n_1$ and $l_2 = l_3 - n_2$. We say that the type of this stack slot is $\text{ptr}(l_3, 0)$.

C.3 Join and Meet of a pointer and a number

The join of a pointer with a number results in an i32(\top), except in the case of constant pointers $\text{ptr}(l_c, n_0)$. For a constant pointer, we know that the base address l_c equates to 0 and so the operation $\text{ptr}(l_c, n_0) \sqcup \text{num}(n_1)$, produces i32($n_0 \sqcup n_1$) as its result. The meet of a pointer and number unequivocally results in a i32(\perp), as per our relation definition shown in Figure ??.

C.4 Join and meet of a i32 with a pointer

If the then branch of a if-then-else expression returned i32(n_1) and the else branch returned a constant pointer $\text{ptr}(l_c, n_0)$, the if-then-else expression would return i32($n_0 \sqcup n_1$). If the else branch returned a pointer with a symbolic address $\text{ptr}(l_2, n_2)$ instead, the if-then-else expression would return i32(\top), since the base address l_1 is unknown. For the meet operation between i32's and ptr's, let us imagine that the type of a stack slot at loop iteration n is i32(n_1) and that at loop iteration $n + 1$, the type of the stack slot is a constant pointer $\text{ptr}(l_c, n_0)$. The type of the stack slot is then both these types and so, $\text{ptr}(l_c, n_0 \sqcap n_1)$. If instead, at loop iteration $n + 1$, the type of the stack slot

1373 is a symbolic pointer $\text{ptr}(l_2, n_2)$, the type of the stack slot would be $\text{ptr}(l_c, n_1) \sqcap \text{ptr}(l_2, n_2)$, which
1374 results in $\text{ptr}(\perp, \perp)$.

1375

1376

1377

1378

1379

1380

1381

1382

1383

1384

1385

1386

1387

1388

1389

1390

1391

1392

1393

1394

1395

1396

1397 C.5 Join and meet of a i32 with a number

1398 If the then branch of a if-then-else expression returned a $i32(n_0)$ and the else branch returned
1399 $\text{num}(n_1)$, the if-then-else expression would return $i32(n_0 \sqcup n_1)$. For the meet operation between
1400 $i32$'s and num 's, let us imagine that the type of a stack slot at loop iteration n is $i32(n_0)$ and that at
1401 loop iteration $n + 1$, the type of the stack slot is a constant pointer $\text{num}(n_1)$. The type of the stack
1402 slot is then both these types and so, $\text{num}(n_0 \sqcap n_1)$.

1403 Note that we never explicitly join or meet locations separate from their offsets. When the meet
1404 operation is performed over two pointers with the same symbolic base address l , the result is a
1405 pointer with that base address l and the join of their offsets. When the meet operation is performed
1406 on two pointers with different symbolic locations, we produce a pointer with a \top location and \top
1407 offset. This is analogous to $\text{ptr}(\top)$, since we never construct $\text{ptr}(\top, n)$. Similarly, for join operations,
1408 we never construct $\text{ptr}(\perp, n)$ and $\text{ptr}(\perp, \perp)$ is analogous to $\text{ptr}(\perp)$. However, in both cases, we do
1409 construct $\text{ptr}(l, \top)$ and $\text{ptr}(l, \perp)$.

1410

1411

1412

1413

1414

1415

1416

1417

1418

1419

1420

1421

1422 D CONSTRAINT GENERATION FOR ALL WEBASSEMBLY INSTRUCTIONS

1423

1424

1425

1426

1427

1428

1429

1430

1431

1432

1433

1434

1435

1436

1437 context $C^\alpha ::= \{\text{func } tf^*, \text{ local } \alpha^*, \text{ global } \alpha^*, \text{ table } n^?, \text{ memory } n^?, \text{ label } (\alpha^*, \varsigma)^*, \text{ return } (\alpha^*, \varsigma)^*\}$

1438

1439 Constraints for Instructions

$S; \varsigma; C^\alpha \vdash e : \alpha^* \rightarrow \alpha^*; S'; \varsigma'$

1440

$$\frac{\alpha \text{ fresh}}{S' = S :: [\alpha \doteq \tau_w(c)]} \text{CONSTANT} \quad \frac{\alpha_1 \in \text{dom}(S) \quad \alpha_2 \text{ fresh}}{S' = S :: [\alpha_1 \doteq \text{unop}(\alpha_2)]} \text{UNARY-OPS}$$

1444

$$\frac{\alpha_1 \alpha_2 \in \text{dom}(S) \quad \alpha_3 \text{ fresh}}{S' = S :: [\alpha_1 \doteq \text{binop}(_, \alpha_2, \alpha_3) \wedge \alpha_2 \doteq \text{binop}(\alpha_1, _, \alpha_3) \wedge \alpha_3 \doteq \text{binop}(\alpha_1, \alpha_2, _)])} \text{BINARY-OPS}$$

1445

$$\frac{\alpha_1 \alpha_2 \in \text{dom}(S) \quad \alpha_3 \text{ fresh}}{S' = S :: [\alpha_1 \doteq \text{testop}(_, \alpha_2, \alpha_3) \wedge \alpha_2 \doteq \text{testop}(\alpha_1, _, \alpha_3) \wedge \alpha_3 \doteq \text{testop}(\alpha_1, \alpha_2, _)])} \text{TEST-OPS}$$

1451

$$\frac{\alpha_1 \alpha_2 \in \text{dom}(S) \quad \alpha_3 \text{ fresh}}{S' = S :: [\alpha_1 \doteq \text{relop}(_, \alpha_2, \alpha_3) \wedge \alpha_2 \doteq \text{relop}(\alpha_1, _, \alpha_3) \wedge \alpha_3 \doteq \text{relop}(\alpha_1, \alpha_2, _)])} \text{REL-OPS}$$

1453

$$\frac{\alpha_1 \in \text{dom}(S) \quad \alpha_2 \text{ fresh}}{S' = S :: [\alpha_1 \doteq \text{cvtop } \tau_w.\text{sx?}(_, \alpha_2) \wedge \alpha_2 \doteq \text{cvtop } \tau_w.\text{sx?}(\alpha_1, _)])} \text{CONVERT-OPS}$$

1458

$$\frac{\alpha_1 \in \text{dom}(S) \quad S' = S :: [\alpha_1 <: \text{ptr} \wedge \alpha_2 \doteq \varsigma[\alpha_1]_{(tf.\text{sx?}, a, o)} \wedge \alpha_2 <: \tau_w]}{S; \varsigma; C^\alpha \vdash \tau_w.\text{load}(tp.\text{sx})? a o : \alpha_1 \rightarrow \alpha_2; S'; \varsigma} \text{LOAD}$$

1461

$$\frac{\alpha_1 \alpha_2 \in \text{dom}(S) \quad S' = S :: [\alpha_1 <: \text{ptr} \wedge \alpha_2 <: \tau_w \wedge \varsigma' \doteq \varsigma[\alpha_1 \mapsto \alpha_2]_{(tp^o, a, o)}]}{S; \varsigma; C^\alpha \vdash \tau_w.\text{store } tp^? a o : \alpha_1 \alpha_2 \rightarrow \epsilon; S'; \varsigma'} \text{STORE}$$

1465

Fig. 10. Constraint Generation Rules.

1466

1468

1469

1470

1471	$tf_w = \tau_w^m \rightarrow \tau_w^n \quad \alpha^m \in \text{dom}(S) \quad \alpha^n \text{ fresh}$
1472	$S; \varsigma; C^\alpha, \text{label}(\alpha^n) \vdash e^* : \alpha^m \rightarrow (\alpha')^n; S'; \varsigma' \quad S'' = S' :: [(\alpha \doteq \alpha')^n]$
1473	$\frac{}{S; \varsigma; C^\alpha \vdash \text{block } tf e^* \text{ end} : \alpha_1^n \dots \alpha_m^n \rightarrow \alpha_1^m \dots \alpha_m^m; S''; \varsigma'} \text{BLOCK}$
1474	
1475	$tf_w = \tau_w^m \rightarrow \tau_w^n \quad \alpha^m \in \text{dom}(S) \quad \alpha^n \text{ fresh}$
1476	$S; \varsigma; C^\alpha, \text{label}(\alpha^m) \vdash e^* : \alpha^m \rightarrow (\alpha')^n; S'; \varsigma' \quad S'' = S' :: [(\alpha \doteq \alpha')^n]$
1477	$\frac{}{S; \varsigma; C^\alpha \vdash \text{loop } tf_w e^* \text{ end} : \alpha^m \rightarrow \alpha^n; S''; \varsigma'} \text{Loop}$
1478	
1479	$tf_w = \tau_w^m \rightarrow \tau_w^n \quad \alpha^m \in \text{dom}(S) \quad \alpha_t^n, \alpha_e^n, \alpha^n \text{ fresh} \quad S' = S :: [\alpha <: \text{num}]$
1480	$S'; \varsigma; C^\alpha, \text{label}(\alpha_t^n) \vdash e_t^* : \alpha_m \rightarrow \alpha_t^n; S_e; \varsigma_e \quad S'; \varsigma; C^\alpha, \text{label}(\alpha_e^n) \vdash e_e^* : \alpha_m \rightarrow \alpha_e^n; S_t; \varsigma_t$
1481	$S'' = S' :: S_e :: S_t :: [\varsigma' \doteq \bigsqcup \varsigma_e \varsigma_t \wedge (\alpha' \doteq \bigsqcup \alpha_t \alpha_e)^n]$
1482	$\frac{}{S; \varsigma; C^\alpha \vdash \text{if } tf e_t^* \text{ else } e_e^* \text{ end} : \alpha^m \alpha \rightarrow \alpha^n; S''; \varsigma'} \text{IF-ELSE}$
1483	
1484	$\alpha^* \alpha^n \in \text{dom}(S) \quad C_{\text{label}}^\alpha(i) = (\alpha')^n, \varsigma' \quad S' = S :: [\varsigma \doteq \bigsqcup \varsigma \varsigma' \wedge (\alpha \doteq \bigsqcup \alpha \alpha')^n]$
1485	$\frac{}{S; \varsigma; C^\alpha \vdash \text{br } i : \alpha^* \alpha_1 \dots \alpha_n \rightarrow \alpha^*; S'; \varsigma} \text{BR}$
1486	
1487	$\alpha^m \alpha \in \text{dom}(S) \quad C_{\text{label}}^\alpha(i) = (\alpha')^n, \varsigma' \quad S' = S :: [\varsigma \doteq \bigsqcup \varsigma \varsigma' \wedge \alpha <: \text{num} \wedge (\alpha \doteq \bigsqcup \alpha \alpha')^n]$
1488	$\frac{}{S; \varsigma; C^\alpha \vdash \text{br_if } i : \alpha^m \alpha \rightarrow \alpha^n; S'; \varsigma} \text{BR-IF}$
1489	
1490	$\alpha^* \alpha^m \alpha \in \text{dom}(S) \quad (C_{\text{label}}^\alpha(i) = (\alpha')^n)^+, \varsigma'$
1491	$S' = S :: ([\varsigma \doteq \bigsqcup \varsigma \varsigma' \wedge \alpha_{n+1} <: \text{num} \wedge \alpha_1 \doteq \bigsqcup \alpha_1 \alpha'_1 \wedge \dots \wedge \alpha_n \doteq \bigsqcup \alpha_n \alpha'_n])^+$
1492	$\frac{}{S; \varsigma; C^\alpha \vdash \text{br_table } i^+ : \alpha^* \alpha_1 \dots \alpha_n \alpha_{n+1} \rightarrow \alpha^*; S'; \varsigma} \text{BR-TABLE}$
1493	
1494	
1495	$\alpha^* \alpha_1 \dots \alpha_n \in \text{dom}(S) \quad C_{\text{return}}(i) = \alpha'_1 \dots \alpha'_n$
1496	$S' = S :: [\alpha_1 \doteq \bigsqcup \alpha_1 \alpha'_1 \wedge \dots \wedge \alpha_n \doteq \bigsqcup \alpha_n \alpha'_n]$
1497	$\frac{}{S; \varsigma; C \vdash \text{return} : \alpha^* \alpha_1 \dots \alpha_n \rightarrow \alpha^*; S'; \varsigma} \text{return}$
1498	
1499	$C_{\text{func}} = tf \quad tf = \tau_w^1 \dots \tau_w^n \rightarrow \tau_w^1 \dots \tau_w^m$
1500	$\alpha_1 \dots \alpha_n \in \text{dom}(S) \quad S' = S :: [\alpha'_1 <: \tau_1 \wedge \dots \wedge \alpha'_m <: \tau_m]$
1501	$\frac{}{S; \varsigma; C \vdash \text{call} : \alpha_1 \dots \alpha_n \rightarrow \alpha'_1 \dots \alpha'_m; S'; \varsigma} \text{call}$
1502	
1503	$C_{\text{table}} = n \quad tf = \tau_w^1 \dots \tau_w^n \rightarrow \tau_w^1 \dots \tau_w^m \quad \alpha_1 \dots \alpha_n \alpha_{n+1} \in \text{dom}(S)$
1504	$S' = S :: [\alpha_{n+1} <: \text{num} \wedge \alpha'_1 <: \tau_1 \wedge \dots \wedge \alpha'_m <: \tau_m]$
1505	$\frac{}{S; \varsigma; C \vdash \text{call_indirect } tf : \alpha_1 \dots \alpha_n \alpha_{n+1} \rightarrow \alpha'_1 \dots \alpha'_m; S'; \varsigma} \text{call_indirect}$
1506	
1507	$\alpha \text{ fresh} \quad S' = S :: [\alpha <: \text{num}]$
1508	$\frac{}{S; \varsigma; C^\alpha \vdash \text{current_memory} : \epsilon \rightarrow \alpha; S'; \varsigma} \text{CURRENT-MEMORY}$
1509	
1510	$\alpha_1 \in \text{dom}(S) \quad \alpha_2 \text{ fresh} \quad S' = S :: [\alpha_1 <: \text{num} \wedge \alpha_2 <: \text{num}]$
1511	$\frac{}{S; \varsigma; C^\alpha \vdash \text{grow_memory} : \alpha_1 \rightarrow \alpha_2; S'; \varsigma} \text{GROW-MEMORY}$
1512	
1513	$\alpha_1 \alpha_2; \alpha_3 \in \text{dom}(S) \quad \alpha_4 \text{ fresh} \quad S' = S :: [\alpha_3 <: \text{num} \wedge \alpha_4 \doteq \bigsqcup \alpha_1 \alpha_2]$
1514	$\frac{}{S; \varsigma; C^\alpha \vdash \text{select} : \alpha_1 \alpha_2 \alpha_3 \rightarrow \alpha_4; S'; \varsigma} \text{SELECT}$
1515	
1516	
1517	
1518	
1519	

Fig. 11. Constraint Generation Rules cont.

$$\begin{array}{c}
\frac{\alpha_1^* \in \text{dom}(S) \quad \alpha_2^* \text{ fresh}}{S; \varsigma; C^\alpha \vdash \text{unreachable} : \alpha_1^* \rightarrow \alpha_2^*; S; \varsigma} \text{UNREACHABLE} \quad \frac{}{S; \varsigma; C \vdash \text{nop} : \epsilon \rightarrow \epsilon; S; \varsigma} \text{NOP} \\
\\
\frac{\alpha \in \text{dom}(S)}{S; \varsigma; C \vdash \text{drop} : \alpha \rightarrow \epsilon; S; \varsigma} \text{DROP} \quad \frac{\alpha \text{ fresh} \quad S' = S :: [\alpha \doteq C_{\text{local}}(i)]}{S; \varsigma; C \vdash \text{get_local } i : \epsilon \rightarrow \alpha; S'; \varsigma} \text{GET-LOCAL} \\
\\
\frac{\alpha \in \text{dom}(S) \quad C_{\text{local}}(i) = \alpha' \quad S' = S :: [\alpha' \doteq \alpha]}{S; \varsigma; C \vdash \text{set_local } i : \alpha \rightarrow \epsilon; S'; \varsigma} \text{SET-LOCAL} \quad \frac{\alpha \in \text{dom}(S) \quad C_{\text{local}}(i) = \alpha' \quad S' = S :: [\alpha' \doteq \alpha]}{S; \varsigma; C \vdash \text{tee_local } i : \alpha \rightarrow \alpha; S'; \varsigma} \text{TEE-LOCAL} \\
\\
\frac{\alpha \text{ fresh} \quad S' = S :: [\alpha \doteq C_{\text{global}}(i)]}{S; \varsigma; C \vdash \text{get_global } i : \epsilon \rightarrow \alpha; S'; \varsigma} \text{GET-GLOBAL} \quad \frac{\alpha \in \text{dom}(S) \quad C_{\text{global}}(i) = \alpha' \quad S' = S :: [\alpha' \doteq \alpha]}{S; \varsigma; C \vdash \text{set_global } i : \alpha \rightarrow \epsilon; S'; \varsigma} \text{SET-GLOBAL} \\
\\
\frac{tf = \tau_w^1 \dots \tau_w^n \rightarrow \tau_w^1 \dots \tau_w^m \quad \tau^f = \alpha_{n_1} \dots \alpha_{n_n} \rightarrow \alpha_{m_1} \dots \alpha_{m_m} \quad \alpha_{n_1} \dots \alpha_{n_n}, \alpha_{m_1} \dots \alpha_{m_m}, \alpha_{l_1} \dots \alpha_{l_l} \text{ fresh}}{S = [\varsigma \doteq \cdot \wedge \alpha_{n_1} <: \tau_w^1 \wedge \dots \wedge \alpha_{n_n} <: \tau_w^n \wedge \alpha_{m_1} <: \tau_w^1 \wedge \dots \wedge \alpha_{m_m} <: \tau_w^m \wedge \alpha_{l_1} <: \tau_w^1 \wedge \dots \wedge \alpha_{l_l} <: \tau_w^l] \quad S; \varsigma; C, \text{local}(\alpha_{n_1}, \dots, \alpha_{n_n}, \alpha_{l_1}, \dots, \alpha_{l_l}), \text{label}(\alpha_{m_1}, \dots, \alpha_{m_m}), \text{return}(\alpha_{m_1}, \dots, \alpha_{m_m}) \vdash e^* : \tau^f; S'; \varsigma'} \text{SET-GLOBAL} \\
\\
[]; \varsigma; C \vdash ex^* \text{ func } tf \text{local } \tau_w^1 \dots \tau_w^l \ e^* : \tau^f; S'; \varsigma' \\
\end{array}$$

Fig. 12. Constraint Generation Rules cont.

E CONSTRAINTS GENERATED FOR AN EXAMPLE WEBASSEMBLY FUNCTION

Consider the WebAssembly function in Figure 13. At the start of the function it is not immediately obvious if either parameter is a pointer. Line 8 has a memory load operation from the addition of the two parameter, so we can infer that one of the parameters must be a pointer – a careful reader will remember that we do not allow the addition of two pointers in our type system. However, we do not know which one is a pointer. It is only at line 17, where the first parameter is multiplied by 42, that we know that the first parameter is a pointer! However, we would now like to reflect this type back up to the function parameter and check that every instruction that uses this parameter does so in accordance with its type. Hence, a forward analysis by itself would be insufficient since it would only propagate type information forwards, when we need the analysis to also propagate types backwards.

F CONSTRAINT SOLVING ALGORITHM

Algorithm 1 Algorithm for solving constraints described in Figure 5

```

1569 Require:  $Constraint = [\alpha \mapsto t^+; \varsigma \mapsto m^+]$ ,  $infl = [\alpha \mapsto \{\alpha^*, \varsigma^*\}; \varsigma \mapsto \{\alpha^*, \varsigma^*\}]$ 
1570 Ensure: The least Solution  $\rho = [\alpha \mapsto \hat{t}_r; \varsigma \mapsto \Sigma]$ 
1571    $Worklist \leftarrow dom(Constraint)$ 
1572   for all  $\alpha \in dom(Constraint)$  do
1573     if  $\exists \alpha \doteqdot \hat{t}_r \in Constraint[\alpha]$  then  $\rho[\alpha] \doteqdot \hat{t}_r$  else  $\rho[\alpha] = \top$ 
1574   for all  $\varsigma \in dom(Constraint)$  do  $\rho[\varsigma] = \cdot$ 
1575   while  $Worklist \neq \emptyset$  do
1576      $v := Worklist.pop()$ 
1577      $Old_v := \rho[v]$ 
1578     if  $v = \alpha$  then  $EVAL_\alpha(Constraint, \rho, \alpha)$ 
1579     if  $v = \varsigma$  then  $EVAL_\varsigma(Constraint, \rho, \varsigma)$ 
1580     if  $Old_v \neq \rho[v]$  then  $Worklist.append(infl[v])$ 
1581   procedure  $EVAL_\alpha(Constraint, \rho, \alpha)$ 
1582      $\hat{t}_r^\alpha := \rho[\alpha]$ 
1583     for all  $t \in Constraint[\alpha]$  do
1584       if  $t = \alpha <: \tau$  then  $\hat{t}_r^\alpha := \hat{t}_r^\alpha \sqcap \tau(\top)$ 
1585       if  $t = \alpha \doteqdot op(\alpha'_1, \dots, \alpha'_n)$  then  $\hat{t}_r^\alpha := \hat{t}_r^\alpha \sqcap eval_{op}(\rho[\alpha'_1], \dots, \rho[\alpha'_n])$ 
1586       if  $t = \alpha \doteqdot \bigsqcup \alpha'_1 \dots \alpha'_m$  then  $\hat{t}_r^\alpha := \hat{t}_r^\alpha \sqcap (\rho[\alpha'_1] \sqcup \dots \sqcup \rho[\alpha'_m])$ 
1587       if  $t = \alpha \doteqdot \varsigma[\alpha']$  then  $\Sigma := \rho[\varsigma] \wedge \hat{t}_r^{\text{ptr}} \doteqdot \rho[\alpha']$ 
1588       if  $\hat{t}_r^{\text{ptr}} = \text{ptr}(l, n) \wedge (l, n) \mapsto \hat{t}_r \in \Sigma$  then  $\hat{t}_r^{(l, n)} := \hat{t}_r$  else  $\hat{t}_r^{(l, n)} := \top$ 
1589       if  $\exists \text{ptr}(\tau, \tau) \mapsto \hat{t}_r \in \Sigma$  then  $\hat{t}_r^{(\tau, \tau)} := \hat{t}_r$  else  $\hat{t}_r^{(\tau, \tau)} := \perp$ 
1590       if  $\exists \text{ptr}(l, \tau) \mapsto \hat{t}_r \in \Sigma$  then  $\hat{t}_r^{(l, \tau)} := \hat{t}_r$  else  $\hat{t}_r^{(l, \tau)} := \perp$ 
1591       if  $\exists \text{ptr}(l', n') \mapsto \hat{t}_r \in \Sigma$  then
1592         if  $\nexists l \neq l' \in Constraint[\alpha']$  then  $\hat{t}_r^{(l', n')} := \hat{t}_r$ 
1593         else if  $\exists l \neq l' \in Constraint[\alpha'] \wedge n' = \top$  then  $\hat{t}_r^{(l', n')} := \hat{t}_r$ 
1594         else  $\hat{t}_r^{(l', n')} := \perp$ 
1595        $\hat{t}_r^\alpha := \hat{t}_r^\alpha \sqcap \hat{t}_r^{(l, n)} \sqcup \hat{t}_r^{(\tau, \tau)} \sqcup \hat{t}_r^{(l, \tau)} \sqcup \hat{t}_r^{(l', n')}$ 
1596      $\rho[\alpha] \doteqdot \hat{t}_r^\alpha$ 
1597   procedure  $EVAL_\varsigma(Constraint, \rho, \varsigma)$ 
1598      $\Sigma \doteqdot \rho[\varsigma]$ 
1599     for all  $m \in Constraint[\varsigma]$  do
1600       if  $m = \varsigma \doteqdot \Sigma'$  then  $\Sigma := \Sigma' \sqcap \Sigma$ 
1601       if  $m = \varsigma \doteqdot \bigsqcup \varsigma'_1 \dots \varsigma'_m$  then  $\Sigma := (\rho[\varsigma'_1] \sqcup \dots \sqcup \rho[\varsigma'_m]) \sqcap \Sigma$ 
1602       if  $m = \varsigma \doteqdot \varsigma'[\alpha_1 \mapsto \alpha_2]$  then  $\Sigma' := \rho[\varsigma'] \wedge \hat{t}_r^{\text{ptr}} := \rho[\alpha_1] \wedge \hat{t}_r^{\text{data}} := \rho[\alpha_2]$ 
1603         if  $\hat{t}_r^{\text{ptr}} = \text{ptr}(l, n)$  then  $\Sigma := \Sigma \sqcap \Sigma'[(l, n) \mapsto \hat{t}_r^{\text{data}}]$  else  $\Sigma := \Sigma \sqcap \Sigma'$ 
1604         if  $\exists \text{ptr}(\tau, \tau) \mapsto \hat{t}'_r \in \Sigma$  then  $\Sigma := \Sigma[(\tau, \tau) \mapsto \hat{t}'_r \sqcup \hat{t}_r^{\text{data}}]$ 
1605         if  $\exists \text{ptr}(l, \tau) \mapsto \hat{t}'_r \in \Sigma$  then  $\Sigma := \Sigma[(l, \tau) \mapsto \hat{t}'_r \sqcup \hat{t}_r^{\text{data}}]$ 
1606         if  $\exists \text{ptr}(l', n') \mapsto \hat{t}'_r \in \Sigma$  then
1607           if  $\nexists l \neq l' \in Constraint[\alpha']$  then  $\Sigma := \Sigma[(l', n') \mapsto \hat{t}'_r \sqcup \hat{t}_r^{\text{data}}]$ 
1608           else if  $\exists l \neq l' \in Constraint[\alpha'] \wedge n' = \top$  then  $\Sigma := \Sigma[(l', n') \mapsto \hat{t}'_r \sqcup \hat{t}_r^{\text{data}}]$ 
1609          $\rho[\varsigma] \doteqdot \Sigma$ 
1610
1611
1612
1613
1614
1615
1616
1617

```

WebAssembly Function	Value Stack	Generated Constraints
1 (func \$foo		
2 (param i32 i32)		$\alpha_0 <: \text{i32} \wedge \alpha_1 <: \text{i32} \wedge \varsigma \doteq []$
3 (result i32)		$\alpha_5 <: \text{i32}$
4 (local i32 i32 i32)		$\alpha_2 <: \text{i32} \wedge \alpha_3 <: \text{i32} \wedge \alpha_4 <: \text{i32}$
5 local.get 0	$[\alpha_0]$	
6 local.get 1	$[\alpha_0\alpha_1]$	
7 i32.add	$[\alpha_6]$	$\alpha_6 \doteq \text{i32.add}(\alpha_0, \alpha_1, _) \wedge \alpha_0 \doteq \text{i32.add}(_, \alpha_1, \alpha_6) \wedge \alpha_1 \doteq \text{i32.add}(\alpha_0, _, \alpha_6)$
8 i32.load	$[\alpha_7]$	$\alpha_6 <: \text{ptr} \wedge \alpha_7 <: \text{i32} \wedge \alpha_7 \doteq \varsigma[\alpha_6]$
9 local.set 2	$[]$	$\alpha_2 \doteq \alpha_7$
10 local.get 1	$[\alpha_1]$	
11 i32.const 4	$[\alpha_1\alpha_8]$	$\alpha_8 \doteq \text{i32}(4)$
12 i32.add	$[\alpha_9]$	$\alpha_9 \doteq \text{i32.add}(\alpha_1, \alpha_8, _) \wedge \alpha_1 \doteq \text{i32.add}(_, \alpha_8, \alpha_9) \wedge \alpha_8 \doteq \text{i32.add}(\alpha_1, _, \alpha_9)$
13 i32.load	$[\alpha_{10}]$	$\alpha_9 <: \text{ptr} \wedge \alpha_{10} <: \text{i32} \wedge \alpha_{10} \doteq \varsigma[\alpha_9]$
14 local.set 3	$[]$	$\alpha_3 \doteq \alpha_{10}$
15 local.get 0	$[\alpha_0]$	
16 i32.const 42	$[\alpha_0\alpha_{11}]$	$\alpha_{11} \doteq \text{i32}(42)$
17 i32.mul	$[\alpha_{12}]$	$\alpha_{12} \doteq \text{i32.mul}(\alpha_0, \alpha_{11}, _) \wedge \alpha_0 \doteq \text{i32.mul}(_, \alpha_{11}, \alpha_{12}) \wedge \alpha_{11} \doteq \text{i32.mul}(\alpha_0, _, \alpha_{12})$
18 local.set 4	$[]$	$\alpha_4 \doteq \alpha_{12}$
19 local.get 2	$[\alpha_2]$	
20 i32.eqz	$[\alpha_{13}]$	$\alpha_{13} \doteq \text{i32.eqz}(\alpha_2)$
21 (if (param) (result i32))		
22 (then		
23 local.get 3)	$[\alpha_3]$	
24 (else		
25 local.get 4)	$[\alpha_4]$	$\alpha_{14} \doteq \sqcup \alpha_3 \alpha_4$
26))	$[\alpha_{14}]$	$\alpha_{14} \doteq \alpha_5$
1643		<i>High-Level Pseudocode for WebAssembly:</i>
1644		int foo (x, y) {
1645		a = *(x + y);
1646		b = *(y + 4);
1647		c = x * 42;
1648		if a == 0 { b }
1649		else { c }
1650		}
1651		
1652		
1653		
1654		
1655		
1656		
1657		
1658		
1659		
1660		
1661		
1662		
1663		
1664		
1665		
1666		

Fig. 13. WebAssembly function with generated constraints.